

# Popcorn Compiler Internals 101: The Gory Details

# Goals

---

- Understand the design & flow of the Popcorn compiler at a high level
- Peel back the covers on LLVM internals & Popcorn-specific modifications
  - Architecture-agnostic middle-end (LLVM intermediate representation)
  - Architecture-specific back-end (target-specific machine code)
- Dig into runtime stack transformation
  - How compiler-generated metadata is stored & used at runtime
  - How migration is invoked on the *source* node
  - How execution resumes on the *destination* node

# What we **won't** cover

---

- Traditional compiler topics
  - The language frontend or language parsing
  - Instruction selection or scheduling
  - Target-agnostic/specific optimizations
  - **Will** cover topics necessary for understanding the Popcorn compiler
- The internals of Popcorn Linux's kernel
- Performance of the system
- Benefits of migration (the hard research stuff)
  - When to migrate (performance, energy, security)
  - Where to migrate


---

# 50,000 foot view of Popcorn Linux & its compiler

# Popcorn Linux

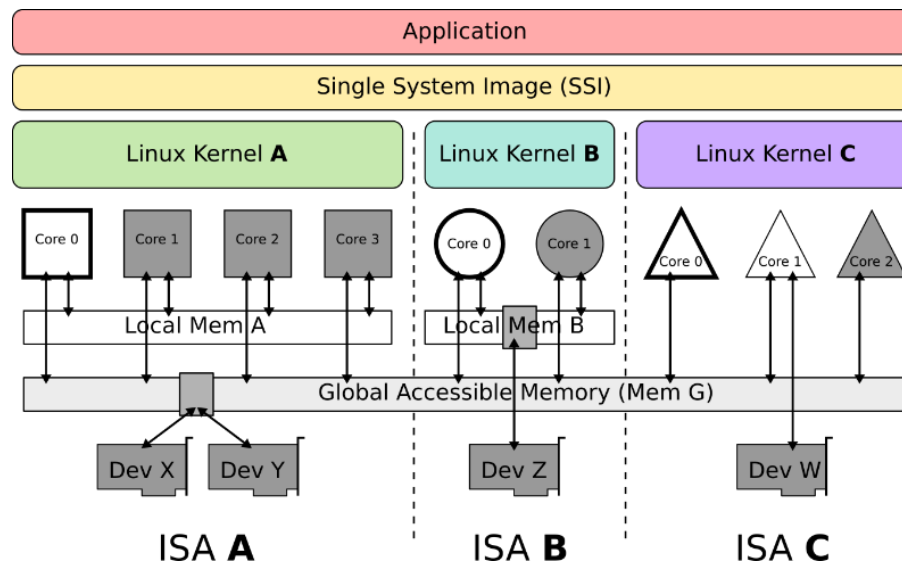
---

- Goal – transparently execute compiled, shared-memory C/C++ applications across physically distinct, heterogeneous-ISA systems
- Developers take advantage of scalability & heterogeneity with no code modifications!



# Popcorn Linux

- Multiple kernels provides *single system image* (SSI) allowing threads to migrate freely between nodes\*
  - Thread migration – stop & resume thread on new node
  - Data migration – transfer data pages between nodes **on-demand**



\*Antonio Barbalace, Robert Lyerly, Christopher Jelesnianski, Anthony Carno, Ho-Ren Chuang, Vincent Legout, and Binoy Ravindran. 2017. Breaking the Boundaries in Heterogeneous-ISA Datacenters. In *Proceedings of the Twenty-Second International Conference on Architectural Support for Programming Languages and Operating Systems* (ASPLOS '17).

# Popcorn Linux

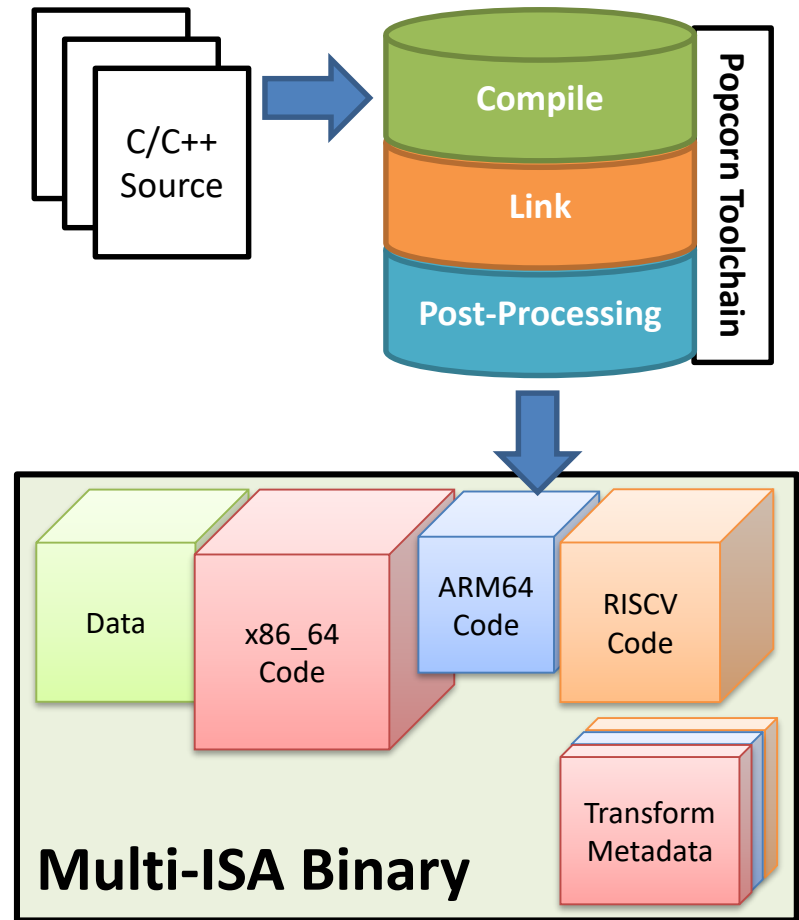
---

- Threads invoke the kernel's migration service through userland syscalls
  - Can't migrate arbitrarily, only at **migration points** (details later)
- Kernels migrate data on-demand through the page-fault mechanism
  - Compiler/runtime need to ensure memory accesses observed by kernel are **semantically equivalent** across all architectures, i.e., pointers reference the same thing and are accessed the same way on all architectures

99% of the compiler implementation is aimed at satisfying this requirement!

# Popcorn Compiler Toolchain


- Compiler toolchain builds **multi-ISA binaries**
  - Create mostly-common virtual address space (data, code, heap)
    - Pointers are valid across all ISAs
  - Dynamically transform thread execution state (stack, registers) between ISA-specific formats at migration time
    - Transform pointers to be valid
  - Instrument generated code with migration points





# Popcorn Compiler Toolchain

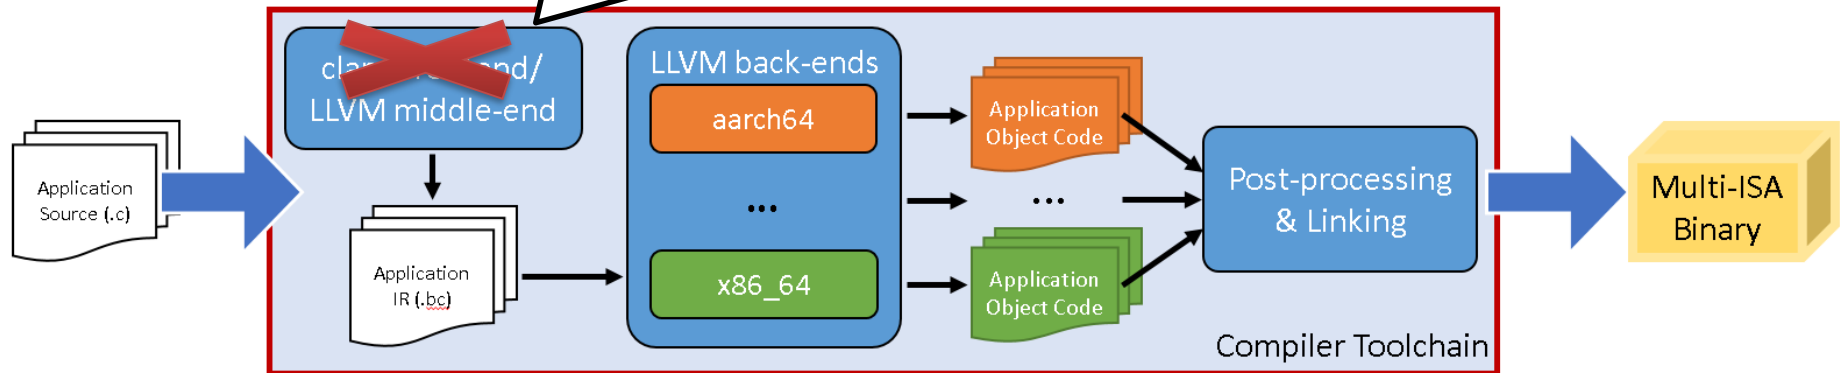
- Built on top of clang/LLVM
  - clang/LLVM 3.7.1, GNU gold 2.27, musl-libc 1.1.18
  - Custom address space alignment, post-processing tools
  - State transformation/migration libraries



# Popcorn Compiler Toolchain

- Built on top of clang/LLVM

- clang
  - Custom
  - State
- We don't need to mess with the frontend – we only need to modify the middle- & back-end
  - Compiler *should* extend to any language that can be lowered to LLVM bytecode, e.g., Rust, Javascript, etc. (YMMV)



# Assumptions

---

- The ISAs have the following characteristics
  - 64-bit address space, meaning pointers are 64 bits
  - Little-endian data format
    - Some RISC architectures can switch between endianness formats, e.g., ARM64 & POWER8
    - Could potentially be relaxed for code, e.g., SPARCV9 allows little-endian data but requires big-endian code
  - Primitive data types have the same sizes & alignments
    - Characters – 8 bits, shorts – 16 bits, integers/longs – 32 bits
    - Long longs/pointers – 64 bits
    - Single precision floating point – 32 bits, double – 64 bits (IEEE format)
- Applications are statically linked, no dynamic libraries

# Assumptions

- The ISAs have the following characteristics
  - 64-bit address space, meaning pointers are 64 bits
  - Little-endian data format
    - Some RISC architectures, e.g., ARM64 & POWER8
    - Could potentially be big-endian, but requires big-endian data but requires big-endian
  - Primitive data types have the same sizes & alignments
    - Characters – 8 bits, shorts – 16 bits, integers/longs – 32 bits
    - Long longs/pointers – 64 bits
    - Single precision floating point – 32 bits, double – 64 bits (IEEE format)
- Applications are statically linked, no dynamic libraries

Forces compiler to lay out all data, i.e., primitives, arrays, structs/objects, in the same format across all architectures

# Limitations

---

- No inline assembly – opaque to liveness analysis in middle-/back-end
- No architecture-specific extensions
  - E.g., crypto instructions, ISA-specific SIMD, etc.
- Note: by limitations, we mean you **cannot have state created by these limitations live during migration – stack transformation can't handle it!**
  - E.g., you cannot have live data in SIMD registers during migration

---

# **Part 1:**

# **A day in the life of an LLVM compilation**

---

# Part 1:

## A day in the life of an LLVM compilation

All source & log files are available  
in the “vanilla-compile” folder

# Generating LLVM IR

---

```
$ clang -O2 -c hello.c  
$ ls  
hello.c  hello.o
```



# Generating LLVM IR

---

What's actually happening on the inside?  
clang/LLVM are designed to be very modular –  
let's break it down...

```
$ clang -O2 -c hello.c  
$ ls  
hello.c  hello.o
```

# Generating LLVM IR

---

hello.c

```
#include <stdio.h>

int main(int argc, char** argv)
{
    printf("Hello, world!\n");
    return 0;
}
```

# Generating LLVM IR

---

```
clang -O2 -emit-llvm -S hello.c
```

hello.c

```
#include <stdio.h>

int main(int argc, char** argv)
{
    printf("Hello, world!\n");
    return 0;
}
```

# Generating LLVM IR

---

```
clang -O2 -emit-llvm -S hello.c
```

hello.c

```
#include <stdio.h>

int main(int argc, char** argv)
{
    printf("Hello, world!\n");
    return 0;
}
```

- -emit-llvm: lower C code to LLVM IR
- -S: emit in human-readable format

# Generating LLVM IR

```
clang -O2 -emit-llvm -S hello.c
```

hello.ll

hello.c

```
#include <stdio.h>

int main(int argc, char** argv)
{
    printf("Hello, world!\n");
    return 0;
}
```



```
; ModuleID = 'hello.c'
target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

@str = private unnamed_addr constant [14 x i8] c"Hello, world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                 i8** nocapture readonly %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* getelementptr inbounds
                                ([14 x i8], [14 x i8]* @str, i64 0, i64 0))
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8* nocapture readonly) #1
```

# Generating LLVM IR

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

@str = private unnamed_addr constant [14 x i8] c"Hello,
world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                 i8** nocapture readnone %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* getelementptr inbounds
                                ([14 x i8], [14 x i8]* @str, i64 0, i64 0))
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8* nocapture readonly) #1
```

- Bitcode is ISA-agnostic, but looks like low-level assembly
  - Conceptually, assembly for some virtual machine or virtual ISA (the “VM” in LLVM)
  - Program objects (functions, global data, etc.) have not been placed in the virtual address space
  - No stack frame particulars (e.g., return address location)
  - No registers or stack slots for program variables (“values” in LLVM)
- See the [language reference manual](#) for a detailed explanation of LLVM’s IR

# Generating LLVM IR

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

@str = private unnamed_addr constant [14 x i8] c"Hello,
world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                 i8** nocapture readnone %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* getelementptr inbounds
                                ([14 x i8], [14 x i8]* @str, i64 0, i64 0))
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8* nocapture readonly) #1
```

Layouts, sizes & alignments of primitive types, e.g., integers, floating-point

# Generating LLVM IR

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

@str = private unnamed_addr constant [14 x i8] c"Hello,
world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                 i8** nocapture readnone %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* getelementptr inbounds
                                ([14 x i8], [14 x i8]* @str, i64 0, i64 0))
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8* nocapture readonly) #1
```

Target architecture/ABI



# Generating LLVM IR

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

@str = private unnamed_addr constant [14 x i8] c"Hello,
world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                 i8** nocapture readnone %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* getelementptr inbounds
                                ([14 x i8], [14 x i8]* @str, i64 0, i64 0))
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8* nocapture readonly) #1
```

Global variable definition

# Generating LLVM IR

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

@str = private unnamed_addr constant [14 x i8] c"Hello,
world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                 i8** nocapture readnone %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* getelementptr inbounds
                                ([14 x i8], [14 x i8]* @str, i64 0, i64 0))
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8* nocapture readonly) #1
```

Function definition

# Generating LLVM IR

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

@str = private unnamed_addr constant [14 x i8] c"Hello,
world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                 i8** nocapture readnone %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* getelementptr inbounds
                                ([14 x i8], [14 x i8]* @str, i64 0, i64 0))
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8* nocapture readonly) #1
```

Function body in single-static assignment form (SSA), organized as a series of basic blocks containing operations

- **Basic block:** a sequence of non-control flow instructions terminated by control flow, e.g., branch or jump

# Generating LLVM IR

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

@str = private unnamed_addr constant [14 x i8] c"Hello,
world!\00"


; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                 i8** nocapture readnone %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* getelementptr inbounds
                                ([14 x i8], [14 x i8]* @str, i64 0, i64 0))
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8* nocapture readonly) #1
```

Function declaration – definition  
lives in another source file or  
“compilation module” in LLVM  
terminology

# Optimizing LLVM IR

- The optimizer (`opt`) is LLVM's “middle-end”
  - Don't worry, `clang` includes these optimizations!
- LLVM applies a large number of target-agnostic optimizations to the IR
  - Optimizations structured as a series of *passes* run over the IR, which consume IR & produce (hopefully) optimized IR



# Optimizing LLVM IR

---

- What is single static assignment?
  - Program variables are lowered to **values** which have the following characteristics:
    - Each value is assigned **exactly once**
    - Each value **must be defined before its use**
    - Once assigned, a value is **immutable**
  - Program variables assigned multiple times are lowered to distinct values (**versioned**) by the compiler
  - LLVM bitcode has no notion of a variable inside of functions – only values!

# Optimizing LLVM IR

---

- Why single static assignment?
  - Provides lots of useful information by construction
    - Explicit use-def chains of a value in a function
    - Liveness ranges for determining when a value is live inside a function
  - Enables many useful analyses & compiler optimizations
    - Instruction scheduling
    - Liveness analysis/register allocation
    - Tons of optimizations – see [Wikipedia](https://en.cppreference.com/w/cpp/string/basic/basic_string_view) for examples

# Lowering IR to machine code

---

- The system compiler (`llc`) is LLVM's back-end
  - Again, `clang` includes this too!
  - Implements the semantics of the IR using operations defined by the target's ISA
- LLVM implements code generation through its [target-independent code generator](#)
  - Another series of passes analyze and transform bitcode to assembly
  - Targets are “plugins” which describe opcodes, registers, ABIs, etc.
  - Most target-specific code lowering is implemented in a target-independent manner!



# Lowering IR to machine code

---

- LLVM has a complex pattern-matching/graph-based framework for instruction selection & scheduling
  - The subject of several Ph.D. theses, not this tutorial
- The backend lowers bitcode into another type of IR, called machine code IR
  - Also in SSA, but very close to the target architecture

# Lowering IR to machine code

---

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = ""
target triple = ""

@str = ... [14 x i8] c"Hello, world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                i8** %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* ...)
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8*) #1
```

# Lowering IR to machine code

---

llc hello.ll

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = ""
target triple = ""

@str = ... [14 x i8] c"Hello, world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                i8** %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* ...)
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8*) #1
```

# Lowering IR to machine code

llc hello.ll

hello.s

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = ""
target triple = ""

@str = ... [14 x i8] c"Hello, world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                i8** %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* ...)
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8*) #1
```



```
.text
.file      "hello.ll"
.globl     main
.align     16, 0x90
.type      main,@function

main:
    .cfi_startproc                                # @main
# BB#0:
    pushq   %rax                                # %entry
.Ltmp0:
    .cfi_def_cfa_offset 16
    movl    $.Lstr, %edi
    callq   puts
    xorl    %eax, %eax
    popq    %rdx
    retq

.Lfunc_end0:
    .size   main, .Lfunc_end0-main
    .cfi_endproc

.type      .Lstr,@object # @str
.section
.rodata.str1.1,"aMS",@progbits,1

.Lstr:
    .asciz  "Hello, world!"
    .size   .Lstr, 14
```

# Lowering IR to machine code

llc hello.ll

hello.s

hello.ll

```
; ModuleID = 'hello.c'
target datalayout = ""
target triple = ""

@str = ... [14 x i8] c"Hello, world!\00"

; Function Attrs: nounwind uwtable
define i32 @main(i32 %argc,
                i8** %argv) #0 {
entry:
    %puts = tail call i32 @puts(i8* ...)
    ret i32 0
}

; Function Attrs: nounwind
declare i32 @puts(i8*) #1
```

But in the  
middle...

```
.text
.file      "hello.ll"
.globl     main
.align     16, 0x90
.type      main,@function

main:
    .cfi_startproc
# BB#0:
    pushq   %rax
.Ltmp0:
    .cfi_def_cfa_offset 16
    movl    $.Lstr, %edi
    callq   puts
    xorl    %eax, %eax
    popq    %rdx
    retq

.Lfunc_end0:
    .size    main, .Lfunc_end0-main
    .cfi_endproc

.type      .Lstr,@object # @str
.section
.rodata.str1.1,"aMS",@progbits,1

.Lstr:
    .asciz   "Hello, world!"
    .size    .Lstr, 14
```

# Lowering IR to machine code

---

```
llc -debug-only=regalloc hello.ll
```

```
***** MACHINEINSTRS *****
```

```
# Machine code for function main: Post SSA
```

```
0B      BB#0: derived from LLVM BB %entry
16B      ADJCALLSTACKDOWN64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
32B      %vreg2<def> = MOV32ri64 <ga:@str>; GR32:%vreg2
48B      %vreg3<def> = SUBREG_TO_REG 0, %vreg2, 4; GR64:%vreg3 GR32:%vreg2
64B      %RDI<def> = COPY %vreg3; GR64:%vreg3
80B      CALL64pcrel32 <ga:@puts>, <regmask>, %RSP<imp-use>, %RDI<imp-use,kill>,
          %RSP<imp-def>, %EAX<imp-def,dead>
96B      ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
112B     %vreg5<def> = MOV32r0 %EFLAGS<imp-def,dead>; GR32:%vreg5
128B     %EAX<def> = COPY %vreg5; GR32:%vreg5
144B     RETQ %EAX<kill>
```

# Lowering IR to machine code

All values in middle-end which could *potentially* be held in a register are lowered to **virtual registers** (vregs)

- LLVM starts by assuming a virtual register set with unlimited registers
- Deciding which values are actually placed in registers and which are spilled to the stack is the purpose of the **register allocator**

```
*****
# Machine code

0B      BB#0: derive
16B      ADJCALLSTACKDOWN64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
32B      %vreg2<def> = MOV32ri64 <ga:@str>; GR32:%vreg2
48B      %vreg3<def> = SUBREG_TO_REG 0, %vreg2, 4; GR64:%vreg3 GR32:%vreg2
64B      %RDI<def> = COPY %vreg3; GR64:%vreg3
80B      CALL64pcrel32 <ga:@puts>, <regmask>, %RSP<imp-use>, %RDI<imp-use,kill>,
          %RSP<imp-def>, %EAX<imp-def,dead>
96B      ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
112B     %vreg5<def> = MOV32r0 %EFLAGS<imp-def,dead>; GR32:%vreg5
128B     %EAX<def> = COPY %vreg5; GR32:%vreg5
144B     RETQ %EAX<kill>
```

# Lowering IR to machine code

```
llc -debug-only=regalloc hello.ll
```

```
***** MACHINEINSTRS *****
```

```
# Machine code
```

LLVM bitcode operations are lowered to target-specific opcodes

```
0B      BB#0: derived from BB#1
16B      ADJCALLSTACKDOWN 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
32B      %vreg2<def> = MOV32ri64 <ga:@str>; GR32:%vreg2
48B      %vreg3<def> = SUBREG_TO_REG 0, %vreg2, 4; GR64:%vreg3 GR32:%vreg2
64B      %RDI<def> = COPY %vreg3; GR64:%vreg3
80B      CALL64pcrel32 <ga:@puts>, <regmask>, %RSP<imp-use>, %RDI<imp-use,kill>,
          %RSP<imp-def>, %EAX<imp-def,dead>
96B      ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
112B     %vreg5<def> = MOV32r0 %EFLAGS<imp-def,dead>; GR32:%vreg5
128B     %EAX<def> = COPY %vreg5; GR32:%vreg5
144B     RETQ %EAX<kill>
```



# Lowering IR to machine code

```
llc -debug-only=regalloc hello.ll
```

Operands are encoded as part of the instruction

- Registers, constants, references to other program objects (here, reference to the string literal “Hello, world!”)

```
***** MACHINEINSTRS
# Machine code for functi

0B      BB#0: derived from IR
16B      ADJCALLSTACKDOWN64 0, 0, 0, <imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
32B      %vreg2<def> = MOV32ri64 <ga:@str>; GR32:%vreg2
48B      %vreg3<def> = SUBREG_TO_REG 0, %vreg2, 4; GR64:%vreg3 GR32:%vreg2
64B      %RDI<def> = COPY %vreg3; GR64:%vreg3
80B      CALL64pcrel32 <ga:@puts>, <regmask>, %RSP<imp-use>, %RDI<imp-use,kill>,
          %RSP<imp-def>, %EAX<imp-def,dead>
96B      ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
112B     %vreg5<def> = MOV32r0 %EFLAGS<imp-def,dead>; GR32:%vreg5
128B     %EAX<def> = COPY %vreg5; GR32:%vreg5
144B     RETQ %EAX<kill>
```

# Lowering IR to machine code

```
llc -debug-only=regalloc hello.ll
```

```
***** MACHINEINSTRS *****
```

```
# Machine code for function main: Post SSA
```

Physical registers appear prior to register allocation due to ABI calling conventions

```
0B
16B
32B
48B
64B
80B
96B
112B
128B
144B

MOV32r164 <ga:@str>; GR32:%vreg2
%vreg5<def> = SUBREG_TO_REG 0, %vreg2, 4; GR64:%vreg3 GR32:%vreg2
%RDI<def> = COPY %vreg3; GR64:%vreg3
CALL64pcrel32 <ga:@puts>, <regmask>, %RSP<imp-use>, %RDI<imp-use,kill>,
    %RSP<imp-def>, %EAX<imp-def,dead>
ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
%vreg5<def> = MOV32r0 %EFLAGS<imp-def,dead>; GR32:%vreg5
%EAX<def> = COPY %vreg5; GR32:%vreg5
RETQ %EAX<kill>
```

# Lowering IR to machine code

---

## Post-register allocation

```
***** MACHINEINSTRS *****
# Machine code for function main: Post SSA

0B      BB#0: derived from LLVM BB %entry
16B      ADJCALLSTACKDOWN64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
64B      %EDI<def,dead> = MOV32ri64 <ga:@str>, %RDI<imp-def>
80B      CALL64pcrel32 <ga:@puts>, <regmask>, %RSP<imp-use>, %RDI<imp-use,kill>,
%RSP<imp-def>, %EAX<imp-def,dead>
96B      ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
128B     %EAX<def> = MOV32r0 %EFLAGS<imp-def,dead>
144B     RETQ %EAX<kill>
```

# Lowering IR to machine code

## Post-register allocation

```
***** MACHINEINSTRS *****
# Machine code for function main: Post SSA

0B      BB#0: derived from LLVM BB %entry
16B      ADJCALLSTACKDOWN64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
64B      %EDI<def,dead> = MOV32ri64 <ga:@str>, %RDI<imp-def>
80B      CALL64pcrel32 <ga:@puts>, <regmask>, %RSP<imp-use>, %RDI<imp-use,kill>,
%RSP<imp-def>, %EAX<imp-def,dead>
96B      ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
128B     %EAX<def> = MOV32r0 %EFLAGS<imp-def,dead>
144B     RETQ %EAX<kill>
```

- Vregs are assigned to physical registers or stack slots based on liveness analysis
  - In this case, LLVM was able to eliminate several vreg definitions by allocating them to the same physical register used in the calling conventions (%edi, %eax)

---

# Compiler Goals Redux: In-Depth

# Goals Redux

---

- Remember, we must ensure pointers are used in a **semantically equivalent** way across all architectures
- Approaches
  - Lay out program objects in a common format and at the same virtual address across all compilations
  - Transform program objects between ISA-specific formats (either statically or at migration time), update references accordingly
- The Popcorn compiler utilizes a mixture of both

# Goals Redux

---

- Popcorn lays out applications in a *mostly-common* format
  - Program objects (global data, functions, the heap) placed at identical addresses for all architectures
  - Stack/register set highly optimized for each ISA – **no common format**
    - ISA defines number and types of registers
    - Compiler tailors stack frame to each ISA based on register allocation results, i.e., compiler spills values to stack that it can't put in registers
- Transform stack & registers between formats at migration time, everything else is aligned at link-time

# Goals Redux

---

- Data objects are equivalent across all architectures
  - Same primitive type sizes & alignments, compiler is forced to lay out higher order types in an identical format
  - Can be placed at identical locations at link-time (details later)
- Code **cannot** be in identical format for different ISAs
  - Like register set, operations/operand format is defined by ISA
  - The manner in which a processor implements a given piece of code is dependent on the operations it supports
  - In other words, *a single piece of code compiled for two different architectures, while semantically identical, executes in different ways*



# Goals Redux

```
void vec_add(const int* a, const int* b, int* c, size_t num) {
    size_t i;
    for(i = 0; i < num; i++)
        c[i] = a[i] + b[i];
}
```

vec\_add\_arm.o: file format elf64-littleaarch64

Disassembly of section .text:

```
0000000000000000 <vec_add>:
 0: b40000e3 cbz x3, 1c <vec_add+0x1c>
 4: b8404408 ldr w8, [x0],#4
 8: b8404429 ldr w9, [x1],#4
 c: 0b080128 add w8, w9, w8
10: b8004448 str w8, [x2],#4
14: d1000463 sub x3, x3, #0x1
18: b5ffff63 cbnz x3, 4 <vec_add+0x4>
1c: d65f03c0 ret
```

vec\_add\_x86.o: file format elf64-x86-64

Disassembly of section .text:

```
0000000000000000 <vec_add>:
 0: 48 85 c9      test    %rcx,%rcx
 3: 74 22         je      27 <vec_add+0x27>
 5: 66 66 2e 0f 1f 84 00 data16  nopw  %cs:0x0(%rax,%rax,1)
 c: 00 00 00 00
10: 8b 06         mov     (%rsi),%eax
12: 03 07         add     (%rdi),%eax
14: 89 02         mov     %eax,(%rdx)
16: 48 83 c7 04   add     $0x4,%rdi
1a: 48 83 c6 04   add     $0x4,%rsi
1e: 48 83 c2 04   add     $0x4,%rdx
22: 48 ff c9     dec     %rcx
25: 75 e9         jne     10 <vec_add+0x10>
27: c3           retq
```

# Goals Redux

```
void vec_add(const int* a, const int* b, int* c, size_t num) {  
    size_t i;  
    for(i = 0; i < num; i++)  
        c[i] = a[i] + b[i];  
}
```

vec\_add\_arm.o: file format elf64-littleaarch64

Disassembly of section .text:

```
0000000000000000 <vec_add>:  
0: b40000e3 cbz x3, 1c <vec_add+0x1c>  
4: b8404408 ldr w8, [x0],#4  
8: b8404429 ldr w9, [x1],#4  
c: 0b080128 add w8, w9, w8  
10: b8004448 str w8, [x2],#4  
14: d1000463 sub x3, x3, #0x1  
18: b5ffff63 cbnz x3, 4 <vec_add+0x4>  
1c: d65f03c0 ret
```

Perform addition  
& update pointers

vec\_add\_x86.o: file format elf64-x86-64

Disassembly of section .text:

```
0000000000000000 <vec_add>:  
0: 48 85 c9 test %rcx,%rcx  
3: 74 22 je 27 <vec_add+0x27>  
5: 66 66 2e 0f 1f 84 00 data16 nopw %cs:0x0(%rax,%rax,1)  
c: 00 00 00 00  
10: 8b 06 mov (%rsi),%eax  
12: 03 07 add (%rdi),%eax  
14: 89 02 mov %eax,(%rdx)  
16: 48 83 c7 04 add $0x4,%rdi  
1a: 48 83 c6 04 add $0x4,%rsi  
1e: 48 83 c2 04 add $0x4,%rdx  
22: 48 ff c9 dec %rcx  
25: 75 e9 jne 10 <vec_add+0x10>  
27: c3 retq
```

# Goals Redux

```
void vec_add(const int* a, const int* b, int* c, size_t num) {
    size_t i;
    for(i = 0; i < num; i++)
        c[i] = a[i] + b[i];
}
```

vec\_add\_arm.o: file format elf64-littleaarch64

Disassembly of section .text:

```
0000000000000000 <vec_add>:
0: b40000e3 cbz x3, 1c <vec_add+0x1c>
4: b8404408 ldr w8, [x0],#4
8: b8404429 ldr w9, [x1],#4
c: 0b080128 add w8, w9, w8
10: b8004448 str w8, [x2],#4
14: d1000463 sub x3, x3, #0x1
18: b5ffff63 cbnz x3, 4 <vec_add+0x4>
1c: d65f03c0 ret
```

Perform addition  
& update pointers

Perform addition

Update pointers

vec\_add\_x86.o: file format elf64-x86-64

Disassembly of section .text:

```
0000000000000000 <vec_add>:
0: 48 85 c9      test    %rcx,%rcx
3: 74 22        je      27 <vec_add+0x27>
5: 66 66 2e 0f 1f 84 00 data16  nopw  %cs:0x0(%rax,%rax,1)
c: 00 00 00 00
10: 8b 06        mov     (%rsi),%eax
12: 03 07        add     (%rdi),%eax
14: 89 02        mov     %eax,(%rdx)
16: 48 83 c7 04   add     $0x4,%rdi
1a: 48 83 c6 04   add     $0x4,%rsi
1e: 48 83 c2 04   add     $0x4,%rdx
22: 48 ff c9     dec     %rcx
25: 75 e9        jne     10 <vec_add+0x10>
27: c3          retq
```

# Goals Redux

```
void vec_add(const int* a, const int* b, int* c, size_t num) {
    size_t i;
    for(i = 0; i < num; i++)
        c[i] = a[i] + b[i];
}
```

vec\_add\_arm.o: file format elf64-littleaarch64

Disassembly of section .text:

```
0000000000000000 <vec_add>:
0: b40000e3 cbz x3, 1c <vec_add+0x1c>
4: b8404408 ldr w8, [x0],#4
8: b8404429 ldr w9, [x1],#4
c: 0b080128 add w8, w9, w8
10: b8004448 str w8, [x2],#4
14: d1000463 sub x3, x3, #0x1
18: b5ffff63 cbnz x3, 4 <vec_add+0x4>
1c: d65f03c0 ret
```

Perform addition  
& update pointers

Perform addition

Update pointers

vec\_add\_x86.o: file format elf64-x86-64

Disassembly of section .text:

```
0000000000000000 <vec_add>:
0: 48 85 c9      test    %rcx,%rcx
3: 74 22         je      27 <vec_add+0x27>
5: 66 66 2e 0f 1f 84 00 data16  nopw  %cs:0x0(%rax,%rax,1)
c: 00 00 00 00

10: 8b 06         mov     (%rsi),%eax
12: 03 07         add     (%rdi),%eax
14: 89 02         mov     %eax,(%rdx)
16: 48 83 c7 04   add     $0x4,%rdi
1a: 48 83 c6 04   add     $0x4,%rsi
1e: 48 83 c2 04   add     $0x4,%rdx
22: 48 ff c9     dec     %rcx
25: 75 e9         jne     10 <vec_add+0x10>
27: c3           retq
```

# Goals Redux

```
void vec_add(const int* a, const int* b, int* c, size_t num) {  
    size_t i;  
    for(i = 0; i < num; i++)  
        c[i] = a[i] + b[i];  
}
```

vec\_add\_arm.o: file format elf64-littleaarch64

Disassembly of section .text:

```
0000000000000000 <vec_add>:  
0: b40000e3 cbz x3, 1c <vec_add+0x1c>  
4: b8404408 ldr w8, [x0],#4  
8: b8404429 ldr w9, [x1],#4  
c: 0b080128 add w8, w9, w8  
10: b8004448 str w8, [x2],#4  
14: d1000463 sub x3, x3, #0x1  
18: b5ffff63 cbnz x3, 4 <vec_add+0x4>  
1c: d65f03c0 ret
```

Perform addition  
& update pointers

Perform addition

Update pointers

Safe for migration

vec\_add\_x86.o: file format elf64-x86-64

Disassembly of section .text:

```
0000000000000000 <vec_add>:  
0: 48 85 c9 test %rcx,%rcx  
3: 74 22 je 27 <vec_add+0x27>  
5: 66 66 2e 0f 1f 84 00 data16 nopw %cs:0x0(%rax,%rax,1)  
c: 00 00 00 00  
10: 8b 06 mov (%rsi),%eax  
12: 03 07 add (%rdi),%eax  
14: 89 02 mov %eax,(%rdx)  
16: 48 83 c7 04 add $0x4,%rdi  
1a: 48 83 c6 04 add $0x4,%rsi  
1e: 48 83 c2 04 add $0x4,%rdx  
22: 48 ff c9 dec %rcx  
25: 75 e9 jne 10 <vec_add+0x10>  
27: c3 retq
```

# Goals Redux

---

- Questions in heterogeneous-ISA execution\*
  1. Given a program address in machine code for one ISA, how do we find the equivalent address in code for another ISA (if there is one)?
  2. If we've established such a mapping, how do we generate a transformation of execution state (registers, stack frames) between ISA-specific formats? How do we ensure such a mapping is feasible?

\*David G. von Bank, Charles M. Shub, and Robert W. Sebesta. 1994. A unified model of pointwise equivalence of procedural computations. *ACM Trans. Program. Lang. Syst.* 16, 6 (November 1994), 1842-1874.

# Goals Redux

- Questions in heterogeneous-ISA execution\*
  1. Given a program address in machine code for one ISA, how do we find the equivalent address in code for another ISA (if there is one)?
  2. If we've established such a mapping, how do we generate a transformation of execution state (registers, stack frames) between ISA-specific formats? How do we ensure such a mapping is feasible?

The compiler programmatically selects a set of all such ***equivalence points*** and inserts call-outs to a migration library. These inserted points are called ***migration points***.

\*David G. von Bank, Charles M. Shub, and Robert W. Sebesta. 1994. A unified model of pointwise equivalence of procedural computations. *ACM Trans. Program. Lang. Syst.* 16, 6 (November 1994), 1842-1874.

---

# **Part 2:**

## **A day in the life of a Popcorn compilation**



---

# Part 2:

## A day in the life of a Popcorn compilation

All source & log files are available in the “het-compile” folder

# Generating LLVM IR

---

fizzbuzz.c

```
#include <stdio.h>

void fizzbuzz(unsigned max)
{
    unsigned i;
    for(i = 0; i < max; i++)
    {
        if((i % 5) == 0 && (i % 3) == 0)
            printf("fizzbuzz\n");
        else if((i % 5) == 0)
            printf("fizz\n");
        else if((i % 3) == 0)
            printf("buzz\n");
    }
}
```

# Generating LLVM IR

---

```
clang -O2 -emit-llvm -S fizzbuzz.c
```

fizzbuzz.c

```
#include <stdio.h>

void fizzbuzz(unsigned max)
{
    unsigned i;
    for(i = 0; i < max; i++)
    {
        if((i % 5) == 0 && (i % 3) == 0)
            printf("fizzbuzz\n");
        else if((i % 5) == 0)
            printf("fizz\n");
        else if((i % 3) == 0)
            printf("buzz\n");
    }
}
```

# Generating LLVM IR


```
clang -O2 -emit-llvm -S fizzbuzz.c
```

fizzbuzz.ll

fizzbuzz.c

```
#include <stdio.h>

void fizzbuzz(unsigned max)
{
    unsigned i;
    for(i = 0; i < max; i++)
    {
        if((i % 5) == 0 && (i % 3) == 0)
            printf("fizzbuzz\n");
        else if((i % 5) == 0)
            printf("fizz\n");
        else if((i % 3) == 0)
            printf("buzz\n");
    }
}
```



```
define void @fizzbuzz(i32 %max) #0 {
entry:
    %cmp.23 = icmp eq i32 %max, 0
    br i1 %cmp.23, label %for.end, label %for.body.preheader

for.body.preheader:    ; preds = %entry
    br label %for.body

for.body:    ; preds = %for.body.preheader, %for.inc
    %i.024 = phi i32 [ %inc, %for.inc ], [ 0, %for.body.preheader ]
    %rem = urem i32 %i.024, 5
    %rem2 = urem i32 %i.024, 3
    %cmp3 = icmp eq i32 %rem2, 0
    ... (if-else statement implementation) ...

for.inc:    ; preds = %if.then, %if.else.8, %if.then.11, %if.then.6
    %inc = add nuw i32 %i.024, 1
    %exitcond = icmp eq i32 %inc, %max
    br i1 %exitcond, label %for.end.loopexit, label %for.body

for.end.loopexit:    ; preds = %for.inc
    br label %for.end

for.end:    ; preds = %for.end.loopexit, %entry
    ret void
}
```

# Inserting Migration Points

---

fizzbuzz.ll

```
define void @fizzbuzz(i32 %max) #0 {
entry:
    %cmp.23 = icmp eq i32 %max, 0
    br i1 %cmp.23, label %for.end,
        label %for.body.preheader

for.body.preheader:
    br label %for.body

for.body:
    %i.024 = phi i32 [ %inc, %for.inc ],
        [ 0, %for.body.preheader ]
    ...(if-else statement implementation)...

for.inc:
    ...(increment loop induction variable)...
    br i1 %exitcond, label %for.end.loopexit,
        label %for.body

for.end.loopexit:
    br label %for.end

for.end:
    ret void
}
```

# Inserting Migration Points

```
opt -select-migration-points -migration-points -S -o fizzbuzz-migpoints.ll fizzbuzz.ll
```

fizzbuzz.ll

```
define void @fizzbuzz(i32 %max) #0 {
entry:
    %cmp.23 = icmp eq i32 %max, 0
    br i1 %cmp.23, label %for.end,
        label %for.body.preheader

for.body.preheader:
    br label %for.body

for.body:
    %i.024 = phi i32 [ %inc, %for.inc ],
                [ 0, %for.body.preheader ]
    ...(if-else statement implementation)...

for.inc:
    ...(increment loop induction variable)...
    br i1 %exitcond, label %for.end.loopexit,
        label %for.body

for.end.loopexit:
    br label %for.end

for.end:
    ret void
}
```

# Inserting Migration Points

```
opt -select-migration-points -migration-points -S -o fizzbuzz-migpoints.ll fizzbuzz.ll
```

fizzbuzz.ll

```
define void @fizzbuzz(i32 %max) #0 {
entry:
    %cmp.23 = icmp eq i32 %max, 0
    br i1 %cmp.23, label %for.end,
        label %for.body.preheader

for.body.preheader:
    br label %for.body

for.body:
    %i.024 = phi i32 [ %inc, %for.inc ],
                [ 0, %for.body.preheader ]
    ...(if-else statement implementation)...

for.inc:
    ...(increment loop induction variable)...
    br i1 %exitcond, label %for.end.loopexit,
        label %for.body

for.end.loopexit:
    br label %for.end

for.end:
    ret void
}
```

opt lets you select individual passes to be run over the IR

- Popcorn passes are patched into LLVM at compiler install time
- Explore Popcorn passes/LLVM mods in <https://github.com/llvm/llvm-project/tree/master/patches/llvm/src>

# Inserting Migration Points

```
opt -select-migration-points -migration-points -S -o fizzbuzz-migpoints.ll fizzbuzz.ll
```

fizzbuzz.ll

```
define void @fizzbuzz(i32 %max) #0 {
entry:
    %cmp.23 = icmp eq i32 %max, 0
    br i1 %cmp.23, label %for.end,
        label %for.body.preheader

for.body.preheader:
    br label %for.body

for.body:
    %i.024 = phi i32 [ %inc, %for.inc ],
        [ 0, %for.body.preheader ]
    ...(if-else statement implementation)...

for.inc:
    ...(increment loop induction variable)...
    br i1 %exitcond, label %for.end.loopexit,
        label %for.body

for.end.loopexit:
    br label %for.end

for.end:
    ret void
}
```



fizzbuzz-migpoints.ll

```
define void @fizzbuzz(i32 %max) #0 {
entry:
    call void @check_migrate(void (i8*)* null, i8* null)
    %cmp.23 = icmp eq i32 %max, 0, !popcorn !2
    br i1 %cmp.23, label %for.end,
        label %for.body.preheader

for.body.preheader:
    br label %for.body

for.body:
    %i.024 = phi i32 [ %inc, %for.inc ],
        [ 0, %for.body.preheader ]
    ...(if-else statement implementation)...

for.inc:
    ...(increment loop induction variable)...
    br i1 %exitcond, label %for.end.loopexit,
        label %for.body

for.end.loopexit:
    br label %for.end

for.end:
    call void @check_migrate(void (i8*)* null, i8* null)
    ret void, !popcorn !2
}
```



# Inserting Migration Points

```
opt -select-migration-points -migration-points -S -o fizzbuzz-migpoints.ll fizzbuzz.ll
```

fizzbuzz.ll

fizzbuzz-migpoints.ll

Compiler inserts migration points at beginning & end of all functions by default

- `check_migrate` defined in `libmigrate.a` and linked in by the compiler
- See `SelectMigrationPoints.cpp` in compiler repo for tuning options

```
for.inc:
  ... (increment loop induction variable) ...
  br i1 %exitcond, label %for.end.loopexit,
    label %for.body

for.end.loopexit:
  br label %for.end

for.end:
  ret void
}
```

```
define void @fizzbuzz(i32 %max) #0 {
entry:
  call void @check_migrate(void (i8*)* null, i8* null)
  %cmp.23 = icmp eq i32 %max, 0, !popcorn !2
  br i1 %cmp.23, label %for.end,
    label %for.body.preheader

for.body.preheader:
  br label %for.body

for.body:
  %i.024 = phi i32 [ %inc, %for.inc ],
    [ 0, %for.body.preheader ]
  ... (if-else statement implementation) ...

for.inc:
  ... (increment loop induction variable) ...
  br i1 %exitcond, label %for.end.loopexit,
    label %for.body

for.end.loopexit:
  br label %for.end

for.end:
  call void @check_migrate(void (i8*)* null, i8* null)
  ret void, !popcorn !2
}
```

# Inserting Stackmaps

---

- Need to accomplish the following:
  - Tag program locations (**all call sites**) with a unique ID
  - Record where live values at that location are stored
- Can't do either at the bitcode level!
  - No code layout – what does a program location mean at the IR level?
  - No storage allocated for live values
- Insert ***stackmap*** intrinsic functions into bitcode
  - Record program and storage locations as back-end lowers bitcode to concrete representation
  - Modified for Popcorn Linux
  - See [Stack maps and patch points in LLVM](#) for more details

# Inserting Stackmaps

- Need to accomplish the following:
  - Tag program locations (**all call sites**) with a unique ID
  - Record where live values at location are stored
- Can't do either at the bitcode level.

Pop quiz – why **all** call sites instead of just migration points?

  - No code layout – what does a program location mean at the IR level?
  - No storage allocated for live values
- Insert ***stackmap*** intrinsic functions into bitcode
  - Record program and storage locations as back-end lowers bitcode to concrete representation
  - Modified for Popcorn Linux
  - See [Stack maps and patch points in LLVM](#) for more details

# Inserting Stackmaps

fizzbuzz-migpoints.ll

```
define void @fizzbuzz(i32 %max) #0 {
entry:
  call void @check_migrate(void (i8*)* null, i8* null)
  %cmp.23 = icmp eq i32 %max, 0, !popcorn !2
  br i1 %cmp.23, label %for.end,
      label %for.body.preheader

for.body.preheader:
  br label %for.body

for.body:
  %i.024 = phi i32 [ %inc, %for.inc ],
               [ 0, %for.body.preheader ]
  ... (if-else statement implementation) ...

for.inc:
  ... (increment loop induction variable) ...
  br i1 %exitcond, label %for.end.loopexit,
      label %for.body

for.end.loopexit:
  br label %for.end

for.end:
  call void @check_migrate(void (i8*)* null, i8* null)
  ret void, !popcorn !2
}
```

# Inserting Stackmaps

```
opt -insert-stackmaps -S -o fizzbuzz-stackmaps.ll fizzbuzz-migpoints.ll
```

fizzbuzz-migpoints.ll

```
define void @fizzbuzz(i32 %max) #0 {
entry:
  call void @check_migrate(void (i8*)* null, i8* null)
  %cmp.23 = icmp eq i32 %max, 0, !popcorn !2
  br i1 %cmp.23, label %for.end,
    label %for.body.preheader

for.body.preheader:
  br label %for.body

for.body:
  %i.024 = phi i32 [ %inc, %for.inc ],
    [ 0, %for.body.preheader ]
  ... (if-else statement implementation) ...

for.inc:
  ... (increment loop induction variable) ...
  br i1 %exitcond, label %for.end.loopexit,
    label %for.body

for.end.loopexit:
  br label %for.end

for.end:
  call void @check_migrate(void (i8*)* null, i8* null)
  ret void, !popcorn !2
}
```

# Inserting Stackmaps

```
opt -insert-stackmaps -S -o fizzbuzz-stackmaps.ll fizzbuzz-migpoints.ll
```

fizzbuzz-migpoints.ll

```
define void @fizzbuzz(i32 %max) #0 {
entry:
  call void @check_migrate(void (i8*)* null, i8* null)
  %cmp.23 = icmp eq i32 %max, 0, !popcorn !2
  br i1 %cmp.23, label %for.end,
    label %for.body.preheader


for.body.preheader:
  br label %for.body

for.body:
  %i.024 = phi i32 [ %inc, %for.inc ],
    [ 0, %for.body.preheader ]
  ... (if-else statement implementation) ...

for.inc:
  ... (increment loop induction variable) ...
  br i1 %exitcond, label %for.end.loopexit,
    label %for.body

for.end.loopexit:
  br label %for.end

for.end:
  call void @check_migrate(void (i8*)* null, i8* null)
  ret void, !popcorn !2
}
```



fizzbuzz-stackmaps.ll

```
define void @fizzbuzz(i32 %max) #0 {
entry:
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 0,
    i32 0, i32 %max)

  %cmp.23 = icmp eq i32 %max, 0, !popcorn !2
  br i1 %cmp.23, label %for.end, label %for.body.preheader

...

if.then:
  %puts22 = tail call i32 @puts(i8* getelementptr inbounds
    ([9 x i8], [9 x i8]* @str.4, i64 0, i64 0))
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 1,
    i32 0, i32 %i.024, i32 %max)

  br label %for.inc

...

for.end:
  ; preds =
  %for.end.loopexit, %entry
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 4,
    i32 0)

  ret void, !popcorn !2
}
```

# Inserting Stackmaps

---

```
call void (i64, i32, ...) @llvm.experimental.stackmap(i64 0, i32 0, i32 %max)
```

# Inserting Stackmaps

---

Unique call-site ID – made unique across all compiled files during post-processing (only have global view at link-time)

```
call void (i64, i32, ...) @llvm.experimental.stackmap(i64 0, i32 0, i32 %max)
```



# Inserting Stackmaps

---

Shadow bytes – unused by Popcorn

```
call void (i64, i32, ...) @llvm.experimental.stackmap(i64 0, i32 0, i32 %max)
```

# Inserting Stackmaps

---

List of live values at this program location

```
call void (i64, i32, ...) @llvm.experimental.stackmap(i64 0, i32 0, i32 %max)
```

# Inserting Stackmaps

---

```
call void (i64, i32, ...) @llvm.experimental.stackmap(i64 0, i32 0, i32 %max)
```

We **create** equivalence points by inserting stackmaps into the bitcode

- Tags equivalent program locations across all compilations, as the **same IR** is lowered through multiple target-specific back-ends
- Lists **all live values** at equivalence point – back-ends are simply responsible for recording where live values are located
- Back-ends must **not** optimize across stackmaps as this violates these invariants (see slide notes)

# Generating Stack Metadata

fizzbuzz-stackmaps.ll

```
define void @fizzbuzz(i32 %max) #0 {
entry:
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 0,
                                                    i32 0, i32 %max)
  %cmp.23 = icmp eq i32 %max, 0, !popcorn !2
  br i1 %cmp.23, label %for.end, label %for.body.preheader

...

if.then:
  %puts22 = tail call i32 @puts(i8* getelementptr inbounds
                                ([9 x i8], [9 x i8]* @str.4, i64 0, i64 0))
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 1,
                                                    i32 0, i32 %i.024, i32 %max)
  br label %for.inc

...

for.end:
                                ; preds =
%for.end.loopexit, %entry
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 4,
                                                    i32 0)
  ret void, !popcorn !2
}
```

# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll
```

```
fizzbuzz-stackmaps.ll
```

```
define void @fizzbuzz(i32 %max) #0 {
entry:
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 0,
                                                    i32 0, i32 %max)
  %cmp.23 = icmp eq i32 %max, 0, !popcorn !2
  br i1 %cmp.23, label %for.end, label %for.body.preheader

...

if.then:
  %puts22 = tail call i32 @puts(i8* getelementptr inbounds
                                ([9 x i8], [9 x i8]* @str.4, i64 0, i64 0))
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 1,
                                                    i32 0, i32 %i.024, i32 %max)
  br label %for.inc

...

for.end:                                     ; preds =
%for.end.loopexit, %entry
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 4,
                                                    i32 0)
  ret void, !popcorn !2
}
```

# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll
```

fizzbuzz-stackmaps

```
define void @fizzbuzz
entry:
  call void @check_mig
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 1,
    i32 0, i32 %i.024, i32 %max)

  %cmp.23 = icmp eq i32 %i.024, %max
  br i1 %cmp.23, label %for.inc, label %for.end

...

if.then:
  %puts22 = tail call i32 @puts(i8* getelementptr inbounds
    ([9 x i8], [9 x i8]* @str.4, i64 0, i64 0))
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 1,
    i32 0, i32 %i.024, i32 %max)
  br label %for.inc

...

for.end:
  ; preds = %for.end.loopexit, %entry
  %for.end.loopexit, %entry
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 4,
    i32 0)
  ret void, !popcorn !2
}
```

- `-mllvm`: Pass option directly to LLVM's middle-/back-end
- `-optimize-regalloc`: use an optimizing register allocator (versus LLVM's fast allocator), which calculates live value ranges in the backend; required by Popcorn's analyses

# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll
```

fizzbuzz-stackmaps.ll


```
define void @fizzbuzz(i32 %max) #0 {
entry:
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 0,
                                                    i32 0, i32 %max)
  %cmp.23 = icmp eq i32 %max, 0, !popcorn !2
  br i1 %cmp.23, label %for.end, label %for.body.preheader

...

if.then:
  %puts22 = tail call i32 @puts(i8* getelementptr inbounds
                                ([9 x i8], [9 x i8]* @str.4, i64 0, i64 0))
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 1,
                                                    i32 0, i32 %i.024, i32 %max)
  br label %for.inc

...

for.end:                                     ; preds =
%for.end.loopexit, %entry
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 4,
                                                    i32 0)
  ret void, !popcorn !2
}
```



# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll
```

```
fizzbuzz-stackmaps.ll
```


```
define void @fizzbuzz(i32 %max) #0 {
entry:
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 0,
    i32 0, i32 %max)
  %cmp.23 = icmp eq i32 %max, 0, !popcorn !2
  br i1 %cmp.23, label %for.end, label %for.body.preheader

...

if.then:
  %puts22 = tail call i32 @puts(i8* getelementptr inbounds
    ([9 x i8], [9 x i8]* @str.4, i64 0, i64 0))
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 1,
    i32 0, i32 %i.024, i32 %max)
  br label %for.inc

...

for.end:                                ; preds =
%for.end.loopexit, %entry
  call void @check_migrate(void (i8*)* null, i8* null)
  call void (i64, i32, ...) @llvm.experimental.stackmap(i64 4,
    i32 0)
  ret void, !popcorn !2
}
```





# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll -mllvm -debug-only=regalloc
```

```
***** MACHINEINSTRS *****
# Machine code for function fizzbuzz: Post SSA
Function Live Ins: %EDI in %vreg4

0B          BB#0: derived from LLVM BB %entry
           Live Ins: %EDI
16B          %vreg5<def> = COPY %EDI; GR32:%vreg5
80B          ADJCALLSTACKDOWN64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
96B          %EDI<def,dead> = MOV32r0 %EFLAGS<imp-def,dead>, %RDI<imp-def>
112B         %ESI<def,dead> = MOV32r0 %EFLAGS<imp-def,dead>, %RSI<imp-def>
128B         CALL64pcrel32 <ga:@check_migrate>, <regmask>, %RSP<imp-use>, %RDI<imp-use>,
                                           %RSI<imp-use,kill>
144B         ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
160B         ADJCALLSTACKDOWN64 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
176B         STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
192B         ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>
208B         CMP32ri8 %vreg5, 0, %EFLAGS<imp-def>; GR32:%vreg5
224B         JE_1 <BB#10>, %EFLAGS<imp-use,kill>
           Successors according to CFG: BB#10 BB#1
```

# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll -mllvm -debug-only=regalloc
```

```
***** MACHINEINSTRS *****  
# Machine code for function fibbuzz: Post SSA  
Function Live Ins: %EDI in BB#0:  
  
0B      BB#0: derived from BB#1  
        Live Ins: %RSP<imp-use>  
16B     %vreg5 = MOVQ $0, %vreg5  
80B     ADJCALLSTACKDOWN64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>  
96B     %EDI = LEAQ -16(%RDI), %EDI  
112B    %ESI = MOVQ %EDI, %ESI  
128B    CALLQ @plt.0, %RDI<imp-use>, %RAX<kill>  
144B    ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>  
160B    ADJCALLSTACKDOWN64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>  
176B    STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5  
192B    ADJCALLSTACKUP64 0, 0, %RSP<imp-def>, %EFLAGS<imp-def,dead>, %RSP<imp-use>  
208B    CMP32ri8 %vreg5, 0, %EFLAGS<imp-def>; GR32:%vreg5  
224B    JE_1 <BB#10>, %EFLAGS<imp-use,kill>  
Successors according to CFG: BB#10 BB#1
```

# Generating Stack Metadata

---

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll -mllvm -debug-only=stacktransform
```

```
***** STACK TRANSFORMATION METADATA *****
***** Function: fizzbuzz
***** REGISTER MAP *****
[%vreg5 -> %R14D] GR32
[%vreg11 -> %EDX] GR32
[%vreg12 -> %AL] GR8
...
```

```
*** Stack slot copies ***
```

```
Stackmap 0:
  STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
```

```
i32 %max: in register %R14D (vreg 5)
```

```
Duplicate operand locations:
```

```
...
*** Finding architecture-specific live values ***
```

```
STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
-> Call instruction SlotIndex 128B, searching vregs 0 -> 31 and stack slots 0 -> 0
```

# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll -mllvm -debug-only=stacktransform
```

```
***** STACK TRANSFORMATION METADATA *****
***** Function: fizzbuzz *****
***** REGISTER MAP *****
[%vreg5 -> %R14D] GR32
[%vreg11 -> %EDX] GR32
[%vreg12 -> %AL] GR8
...

*** Stack slot copies ***

Stackmap 0:
  STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5

i32 %max: in register %R14D (vreg 5)

Duplicate operand locations:

...
*** Finding architecture-specific live values ***
```

```
STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
-> Call instruction SlotIndex 128B, searching vregs 0 -> 31 and stack slots 0 -> 0
```

- Metadata emitted by vanilla stackmaps is not complete
  - Intended to allow capturing enough live state to jump to optimized implementation (e.g., hot-patching optimized code in virtual machine)
  - Need to augment with complete frame information
  - See `StackTransformMetadata.cpp` for more details

# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll -mllvm -debug-only=stacktransform
```

```
***** STACK TRANSFORMATION METADATA *****  
***** Function: fizzbuzz *****  
***** REGISTER MAP *****  
[%vreg5 -> %R14D] GR32  
[%vreg11 -> %EDX] GR32  
[%vreg12 -> %AL] GR8  
...
```

Register allocation results

```
*** Stack slot copies ***
```

```
Stackmap 0:  
  STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
```

```
i32 %max: in register %R14D (vreg 5)
```

```
Duplicate operand locations:
```

```
...  
*** Finding architecture-specific live values ***
```

```
STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5  
-> Call instruction SlotIndex 128B, searching vregs 0 -> 31 and stack slots 0 -> 0
```

# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll -mllvm -debug-only=stacktransform
```

```
***** STACK TRANSFORMATION METADATA *****
***** Function: fizzbuzz *****
***** REGISTER MAP *****
[%vreg5 -> %R14D] GR32
[%vreg11 -> %EDX] GR32
[%vreg12 -> %AL] GR8
...
```

```
*** Stack slot copies ***
```

```
Stackmap 0:
  STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
```

```
i32 %max: in register %R14D (vreg 5)
```

```
Duplicate operand locations:
```

```
...
*** Finding architecture-specific live values ***
```

```
STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
```

```
-> Call instruction SlotIndex 128B, searching vregs 0 -> 31 and stack slots 0 -> 0
```

Mapping of IR values -> machine-code values

- Value *may* be live in multiple locations

# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll -mllvm -debug-only=stacktransform
```

```
***** STACK TRANSFORMATION METADATA *****
***** Function: fizzbuzz *****
***** REGISTER MAP *****
[%vreg5 -> %R14D] GR32
[%vreg11 -> %EDX] GR32
[%vreg12 -> %AL] GR8
...
*** Stack slot copies ***

Stackmap 0:
  STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5

i32 %max: in register %R14D (vreg 5)

Duplicate operand locations:
...
*** Finding architecture-specific live values ***
```

Handle architecture-specific live values that may arise due to backend-optimizations or ABI

```
STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
-> Call instruction SlotIndex 128B, searching vregs 0 -> 31 and stack slots 0 -> 0
```

# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll -mllvm -debug-only=stacktransform
```

```
***** STACK TRANSFORMATION METADATA *****
***** Function: fizzbuzz
***** REGISTER MAP *****
[%vreg5 -> %R14D] GR32
[%vreg11 -> %EDX] GR32
[%vreg12 -> %AL] GR8
...
```

```
*** Stack slot copies ***
```

```
Stackmap 0:
  STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
```

```
i32 %max: in register %R14D (vreg 5)
```

```
Duplicate operand locations:
```

```
...
*** Finding architecture-specific li
```

SlotIndex: instruction location in machine code IR

```
STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
```

```
-> Call instruction SlotIndex 128B, searching vregs 0 -> 31 and stack slots 0 -> 0
```



# Generating Stack Metadata

```
clang -c -mllvm -optimize-regalloc -o fizzbuzz.o fizzbuzz-stackmaps.ll -mllvm -debug-only=stacktransform
```

```
***** STACK TRANSFORMATION METADATA *****
***** Function: fizzbuzz
***** REGISTER MAP *****
[%vreg5 -> %R14D] GR32
[%vreg11 -> %EDX] GR32
[%vreg12 -> %AL] GR8
...
```

```
*** Stack slot copies ***
```

```
Stackmap 0:
  STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
```

```
i32 %max: in register
```

```
Duplicate operand
```

```
...
```

```
*** Finding architecture
```

Looking for registers/stack slots live across stackmap but not contained in stackmap

- Back-end will issue warning if it finds live value it can't handle during transformation – **pay attention to these warnings!**

```
STACKMAP 0, 0, %vreg5, ...; GR32:%vreg5
```

```
-> Call instruction SlotIndex 128B, searching vregs 0 -> 31 and stack slots 0 -> 0
```

# Generating Stack Metadata

```
readelf -SW fizzbuzz.o
```

There are 17 section headers, starting at offset 0x5e0:

Section Headers:

[Nr]	Name	Type	Address	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	0000000000000000	000000	000000	00		0	0	0
[ 1]	.strtab	STRTAB	0000000000000000	0004b0	00012f	00		0	0	1
[ 2]	.text	PROGBITS	0000000000000000	000040	0000b9	00	AX	0	0	16
[ 3]	.rela.text	RELA	0000000000000000	0003a8	0000c0	18		16	2	8
[ 4]	.data	PROGBITS	0000000000000000	0000fc	000000	00	WA	0	0	4
[ 5]	.bss	NOBITS	0000000000000000	0000fc	000000	00	WA	0	0	4
[ 6]	.rodata.str1.1	PROGBITS	0000000000000000	0000fc	000013	01	AMS	0	0	1
[ 7]	.comment	PROGBITS	0000000000000000	00010f	000035	01	MS	0	0	1
[ 8]	.note.GNU-stack	PROGBITS	0000000000000000	000144	000000	00		0	0	1
[ 9]	.stack_transform.unwind	PROGBITS	0000000000000000	000144	000010	04			0	0 4
[10]	.stack_transform.unwind_arange	PROGBITS	0000000000000000	000158	000010	10			0	0 8
[11]	.rela.stack_transform.unwind_arange	RELA	0000000000000000	000468	000018	18			16	10 8
[12]	.llvm_stackmaps	PROGBITS	0000000000000000	000168	000118	00	A	0	0	8
[13]	.rela.llvm_stackmaps	RELA	0000000000000000	000480	000018	18			16	12 8
[14]	.eh_frame	PROGBITS	0000000000000000	000280	000038	00	A	0	0	8
[15]	.rela.eh_frame	RELA	0000000000000000	000498	000018	18			16	14 8
[16]	.symtab	SYMTAB	0000000000000000	0002b8	0000f0	18			1	7 8

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings), l (large)  
I (info), L (link order), G (group), T (TLS), E (exclude), x (unknown)  
O (extra OS processing required) o (OS specific), p (processor specific)

# Generating Stack Metadata

```
readelf -SW fizzbuzz.o
```

There are 17 section headers, starting at offset 0x5e0:

Section Headers:

[Nr]	Name	Type	Address	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	0000000000000000	000000	000000	00		0	0	0
[ 1]	.strtab	STRTAB	0000000000000000	000000	000000	00		0	0	1
[ 2]	.text	PROGBITS	0000000000000000	000000	000000	00		0	0	16
[ 3]	.rela.text	RELA	0000000000000000	000000	000000	00		2	8	
[ 4]	.data	PROGBITS	0000000000000000	000000	000000	00		0	4	
[ 5]	.bss	NOBITS	0000000000000000	000000	000000	00		0	4	
[ 6]	.rodata.str1.1	PROGBITS	0000000000000000	000000	000000	00		0	1	
[ 7]	.comment	PROGBITS	0000000000000000	00010f	000035	01	MS	0	0	1
[ 8]	.note.GNU-stack	PROGBITS	0000000000000000	000144	000000	00		0	0	1
[ 9]	.stack_transform.unwind	PROGBITS	0000000000000000	000144	000010	04		0	0	4
[10]	.stack_transform.unwind arange	PROGBITS	0000000000000000	000158	000010	10		0	0	8
[11]	.rela.stack_transform.unwind arange	RELA	0000000000000000	000468	000018	18		16	10	8
[12]	.llvm_stackmaps	PROGBITS	0000000000000000	000168	000118	00	A	0	0	8
[13]	.rela.llvm_stackmaps	RELA	0000000000000000	000480	000018	18		16	12	8
[14]	.eh_frame	PROGBITS	0000000000000000	000280	000038	00	A	0	0	8
[15]	.rela.eh_frame	RELA	0000000000000000	000498	000018	18		16	14	8
[16]	.symtab	SYMTAB	0000000000000000	0002b8	0000f0	18		1	7	8

Metadata describing callee-save information for functions

- Unwind frames from the stack

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings), l (large)  
I (info), L (link order), G (group), T (TLS), E (exclude), x (unknown)  
O (extra OS processing required) o (OS specific), p (processor specific)

# Generating Stack Metadata

```
readelf -SW fizzbuzz.o
```

There are 17 section headers, starting at offset 0x5e0:

Section Headers:

[Nr]	Name	Type	Address	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	0000000000000000	000000	000000	00		0	0	0
[ 1]	.strtab	STRTAB	0000000000000000	0004b0	00012f	00		0	0	1
[ 2]	.text	PROGBITS	0000000000000000	000040	0000b9	00	AX	0	0	16
[ 3]	.rela.text	RELA	0000000000000000	0003a8	0000c0	18		16	2	8
[ 4]	.data	PROGBITS							0	4
[ 5]	.bss	NOBITS							0	4
[ 6]	.rodata.str1.1	PROGBITS							0	1
[ 7]	.comment	PROGBITS							0	1
[ 8]	.note.GNU-stack	PROGBITS							0	1
[ 9]	.stack_transform.unwind_info	PROGBITS							0	0 4
[10]	.stack_transform.unwind_range	PROGBITS	0000000000000000	000158	000010	10			0	0 8
[11]	.rela.stack_transform.unwind_range	RELA		0000000000000000	000468	000018	18		16	10 8
[12]	.llvm_stackmaps	PROGBITS	0000000000000000	000168	000118	00	A	0	0	8
[13]	.rela.llvm_stackmaps	RELA		0000000000000000	000480	000018	18		16	12 8
[14]	.eh_frame	PROGBITS	0000000000000000	000280	000038	00	A	0	0	8
[15]	.rela.eh_frame	RELA		0000498	000018	18		16	14	8
[16]	.symtab	SYMTAB	0000000000000000	0002b8	0000f0	18		1	7	8

Stack transformation metadata

- Program locations (function + offset)
- Live value locations

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings), l (large)  
I (info), L (link order), G (group), T (TLS), E (exclude), x (unknown)  
O (extra OS processing required) o (OS specific), p (processor specific)

# Generating Stack Metadata

---

```
dump-llvm-stackmap -f fizzbuzz.o
```

```
Reading section .llvm_stackmaps: Found 1 stackmaps
Stackmap v1: 1 functions, 0 constants, 5 call sites
  Function 0: address=0, stack size=24, number of unwinding entries: 4, offset into unwinding
section: 0
    Call site 0: function 0, offset @ 19, 1 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 14
    Call site 1: function 0, offset @ 92, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
    Call site 2: function 0, offset @ 121, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
    Call site 3: function 0, offset @ 154, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
    Call site 4: function 0, offset @ 180, 0 locations, 0 live-outs, 0 arch-specific locations
```

# Generating Stack Metadata

```
dump-llvm-stackmap -f fizzbuzz.o
```

```
Reading section .llvm_stackmaps: Found 1 stackmaps
Stackmap v1: 1 functions, 0 constants, 5 call sites
  Function 0: address=0, stack size=24, number of unwinding entries: 4, offset into unwinding
section: 0
    Call site 0: function 0, offset @ 121, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
    Call site 1: function 0, offset @ 154, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
    Call site 2: function 0, offset @ 180, 0 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
    Call site 3: function 0, offset @ 121, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
    Call site 4: function 0, offset @ 154, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
```

Number of functions for which we have metadata,  
number of stackmap call sites across all functions

# Generating Stack Metadata

```
dump-llvm-stackmap -f fizzbuzz.o
```

```
Reading section .llvm_stackmaps: Found 1 stackmaps
Stackmap v1: 1 functions, 0 constants, 5 call sites
  Function 0: address=0, stack size=24, number of unwinding entries: 4, offset into unwinding
section: 0
    Call site 0: function 0, offset @ 19, 1 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
    Call site 1: function 0, offset @ 121, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
    Call site 2: function 0, offset @ 154, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
    Call site 3: function 0, offset @ 180, 0 locations, 0 live-outs, 0 arch-specific locations
```

Per-function metadata describing frame size and how to unwind the frame to the caller

# Generating Stack Metadata

```
dump-llvm-stackmap -f fizzbuzz.o
```

```
Reading section .llvm_stackmaps: Found 1 stackmaps
Stackmap v1: 1 functions, 0 constants, 5 call sites
  Function 0: address=0, stack size=24, number of unwinding entries: 4, offset into unwinding
section: 0
    Call site 0: function 0, offset @ 19, 1 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 14
    Call site 1: function 0, offset @ 92, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 14
      Location: in register 14
    Call site 2: function 0, offset @ 125, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 14
      Location: in register 14
    Call site 3: function 0, offset @ 154, 2 locations, 0 live-outs, 0 arch-specific locations
      Location: in register 3
      Location: in register 14
    Call site 4: function 0, offset @ 180, 0 locations, 0 live-outs, 0 arch-specific locations
```

Per-call site information describing program location & live values (one per stackmap inserted into bitcode)



# Putting It All Together

---

```
$ clang -O2 -popcorn-migratable -c fizzbuzz.c
$ ls
fizzbuzz.c  fuzzbuzz.o  fizzbuzz_x86_64.o
$ file fuzzbuzz.o
fizzbuzz.o: ELF 64-bit LSB relocatable, ARM aarch64, version 1 (GNU/Linux), not stripped
$ file fizzbuzz_x86_64.o
fizzbuzz_x86_64.o: ELF 64-bit LSB relocatable, x86-64, version 1 (GNU/Linux), not stripped
```

# Putting It All Together

---

- Insert migration library call-outs, collect stack transformation metadata at equivalence points
- Generate object files for **all** supported architectures simultaneously (requires `-c`)

```
$ clang -O2 -popcorn-migratable -c fizzbuzz.c
$ ls
fizzbuzz.c  fuzzbuzz.o  fizzbuzz_x86_64.o
$ file fuzzbuzz.o
fizzbuzz.o: ELF 64-bit LSB relocatable, ARM aarch64, version 1 (GNU/Linux), not stripped
$ file fizzbuzz_x86_64.o
fizzbuzz_x86_64.o: ELF 64-bit LSB relocatable, x86-64, version 1 (GNU/Linux), not stripped
```

# Compilation Recap

---

- Inserts call-outs to migration library
- Constructs equivalence points by inserting stackmaps into LLVM bitcode
  - Tags program locations across compilations for all targets
  - Captures locations of all live values at equivalence points
- Generates single set of optimized LLVM bitcode and lowers it through multiple target-specific backends

---

# **Part 3:**

## **Linking & post-processing a Popcorn compilation**

---

# Part 3:

## Linking & post-processing a Popcorn compilation

All source & log files are available  
in the “het-link” folder

# Link-Time Alignment

---

- The compiler has taken care of generating metadata required for stack/register transformation
- Still need to align global objects
  - Statically-allocated global data
  - Code, i.e., functions
  - We **do not** need to worry about dynamically-allocated data (heap)
    - Memory allocator (e.g., `malloc`) is responsible for arranging data in heap
    - Use semantically equivalent memory allocator (musl) on all architectures

# Link-Time Alignment

---

- Linux uses the ***Executable and Linkable Format*** (ELF)
  - Data sections (statically allocated):
    - `.rodata` – read-only global data initialized within the program
    - `.data` – readable/writable global data initialized within the program
    - `.bss` – readable/writable global data **not** initialized within the program
      - Linux initializes to zero
  - Code section:
    - `.text` – ISA-specific machine code generated by the compiler
  - Miscellaneous – symbol/string tables, constructors/destructors , debugging information
  - Popcorn's compiler adds metadata sections required for stack transformation

# Link-Time Alignment

---

- Use *linker scripts* to align symbols across all compilations
  - Program objects referenced by symbol – requires **all** program objects have a symbol attached, including string literals
    - See `NameStringLiterals.cpp` in repo for details
  - After generating “vanilla” (read: unaligned) version of binary for each architecture, generate linker script & re-link
  - See the [documentation on linker scripts](#) for more details



# Link-Time Alignment

---

- Linker scripts can't control placement of symbols, only ELF sections
  - Luckily, ELF format permits arbitrary numbers of sections in object files
  - Solution: place each program object in its own section
    - Requires `-ffunction-sections` & `-fdata-sections` (included automatically with `-popcorn-migratable`)

# Link-Time Alignment

---

fizzbuzz.c

```
#include <stdio.h>

void fizzbuzz(unsigned max)
{
    unsigned i;
    for(i = 0; i < max; i++)
    {
        if((i % 5) == 0 && (i % 3) == 0)
            printf("fizzbuzz\n");
        else if((i % 5) == 0)
            printf("fizz\n");
        else if((i % 3) == 0)
            printf("buzz\n");
    }
}
```

# Link-Time Alignment

---

```
clang -O2 -popcorn-migratable -c fizzbuzz.c
```

fizzbuzz.c

```
#include <stdio.h>

void fizzbuzz(unsigned max)
{
    unsigned i;
    for(i = 0; i < max; i++)
    {
        if((i % 5) == 0 && (i % 3) == 0)
            printf("fizzbuzz\n");
        else if((i % 5) == 0)
            printf("fizz\n");
        else if((i % 3) == 0)
            printf("buzz\n");
    }
}
```


# Link-Time Alignment

```
clang -O2 -popcorn-migratable -c fizzbuzz.c
```

fizzbuzz.c

```
#include <stdio.h>

void fizzbuzz(unsigned max)
{
    unsigned i;
    for(i = 0; i < max; i++)
    {
        if((i % 5) == 0 && (i % 3) == 0)
            printf("fizzbuzz\n");
        else if((i % 5) == 0)
            printf("fizz\n");
        else if((i % 3) == 0)
            printf("buzz\n");
    }
}
```



fizzbuzz.o  
(AArch64)



fizzbuzz\_x86\_64.o  
(x86-64)

# Link-Time Alignment

```
readelf -SW fizzbuzz.o
```

fizzbuzz.o  
(AArch64)

There are 20 section headers, starting at offset 0x7f0:

Section Headers:

[Nr]	Name	Type	Address	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	0000000000000000	000000	000000	00		0	0	0
[ 1]	.strtab	STRTAB	0000000000000000	000640	0001a9	00		0	0	1
[ 2]	.text	PROGBITS	0000000000000000	000040	000000	00	AX	0	0	4
[ 3]	.data	PROGBITS	0000000000000000	000040	000000	00	WA	0	0	4
[ 4]	.bss	NOBITS	0000000000000000	000040	000000	00	WA	0	0	4
[ 5]	.text.fizzbuzz	PROGBITS	0000000000000000	000040	0000fc	00	AX	0	0	4
[ 6]	.rela.text.fizzbuzz	RELA	0000000000000000	0004f0	000108	18		19	5	8
[ 7]	.rodata.fizzbuzz__str_fizzbuzz__	PROGBITS	0000000000000000	00013c	00000a	00	A	0	0	1
[ 8]	.rodata.fizzbuzz__str_1_fizz__	PROGBITS	0000000000000000	000146	000006	00	A	0	0	1
[ 9]	.rodata.fizzbuzz__str_2_buzz__	PROGBITS	0000000000000000	00014c	000006	00	A	0	0	1
[10]	.comment	PROGBITS	0000000000000000	000152	000035	01	MS	0	0	1
[11]	.note.GNU-stack	PROGBITS	0000000000000000	000187	000000	00		0	0	1
[12]	.stack_transform.unwind	PROGBITS	0000000000000000	000188	000010	04		0	0	4
[13]	.stack_transform.unwind_arange	PROGBITS	0000000000000000	000198	000010	10		0	0	8
[14]	.rela.stack_transform.unwind_arange	RELA	0000000000000000	0005f8	000018	18		19	13	8
[15]	.llvm_stackmaps	PROGBITS	0000000000000000	0001a8	000118	00	A	0	0	8
[16]	.rela.llvm_stackmaps	RELA	0000000000000000	000610	000018	18		19	15	8
[17]	.eh_frame	PROGBITS	0000000000000000	0002c0	000038	00	A	0	0	8
[18]	.rela.eh_frame	RELA	0000000000000000	000628	000018	18		19	17	8
[19]	.symtab	SYMTAB	0000000000000000	0002f8	0001f8	18		1	15	8

# Link-Time Alignment

```
readelf -SW fizzbuzz.o
```

fizzbuzz.o  
(AArch64)

There are 20 section headers, starting at offset 0x100000000:

Section Headers:

[Nr]	Name	Type	Address	Size	Align	Flags	Other
[ 0]		NULL					
[ 1]	.strtab	STRTAB	0000000000000000	00000000	00	AX	0 0 4
[ 2]	.text	PROGBITS	0000000000000000	000040	0000fc	00	AX 0 0 4
[ 3]	.data	PROGBITS	0000000000000000	000040	0000fc	00	AX 0 0 4
[ 4]	.bss	NOBITS	0000000000000000	000040	0000fc	00	AX 0 0 4
[ 5]	.text.fizzbuzz	PROGBITS	0000000000000000	000040	0000fc	00	AX 0 0 4
[ 6]	.rela.text.fizzbuzz	RELA	0000000000000000	000040	0000fc	00	AX 0 0 4
[ 7]	.rodata.fizzbuzz__str_fizzbuzz__	PROGBITS	0000000000000000	000013c	000000a	00	A 0 0 1
[ 8]	.rodata.fizzbuzz__str_1_fizz__	PROGBITS	0000000000000000	0000146	0000006	00	A 0 0 1
[ 9]	.rodata.fizzbuzz__str_2_buzz__	PROGBITS	0000000000000000	000014c	0000006	00	A 0 0 1
[10]	.comment	PROGBITS	0000000000000000	0000152	0000035	01	MS 0 0 1
[11]	.note.GNU-stack	PROGBITS	0000000000000000	0000187	0000000	00	0 0 1
[12]	.stack_transform.unwind	PROGBITS	0000000000000000	0000188	0000010	04	0 0 4
[13]	.stack_transform.unwind_range	PROGBITS	0000000000000000	0000198	0000010	10	0 0 8
[14]	.rela.stack_transform.unwind_range	RELA	0000000000000000	00005f8	0000018	18	19 13 8
[15]	.llvm_stackmaps	PROGBITS	0000000000000000	00001a8	0000118	00	A 0 0 8
[16]	.rela.llvm_stackmaps	RELA	0000000000000000	0000610	0000018	18	19 15 8
[17]	.eh_frame	PROGBITS	0000000000000000	00002c0	0000038	00	A 0 0 8
[18]	.rela.eh_frame	RELA	0000000000000000	0000628	0000018	18	19 17 8
[19]	.symtab	SYMTAB	0000000000000000	00002f8	00001f8	18	1 15 8

Each program object is placed into its own section & prepended with the section name corresponding to that object type (e.g., .text for code)

# Link-Time Alignment

---

- We use a modified linker based on GNU's `gold` and a python tool `pyalign` for alignment
- First linking pass
  - Link vanilla (unaligned) version of binary for each target
  - Dump section names, including sizes & alignments, into a ***map file***
- Alignment
  - Parse binaries/map file and generate linker script for each target
- Second linking pass
  - Link heterogeneous (aligned) version of binary for each target using linker scripts generated by `pyalign`

# Link-Time Alignment

---

```
$ ld.gold -L/usr/lib/gcc-cross/aarch64-linux-gnu/5 -Map map_aarch64.txt \  
main.o fizzbuzz.o -o fizzbuzz_aarch64 \  
<install>/aarch64/lib/crt1.o <install>/aarch64/lib/libc.a \  
<install>/aarch64/lib/libmigrate.a <install>/aarch64/lib/libstack-transform.a \  
<install>/aarch64/lib/libelf.a <install>/aarch64/lib/libc.a \  
--start-group -lgcc -lgcc_eh -end-goup
```



# Link-Time Alignment

---

```
$ ld.gold -L/usr/lib/gcc-cross/aarch64-linux-gnu/5 -Map map_aarch64.txt \  
main.o fizzbuzz.o -o fizzbuzz_aarch64 \  
<install>/aarch64/lib/crt1.o <install>/aarch64/lib/crti.o <install>/aarch64/lib/libc.a \  
<install>/aarch64/lib/libmigrate.a <install>/aarch64/lib/libstack-transform.a \  
<install>/aarch64/lib/libstack-transform.a \  
--sta
```

Generate map file named map\_aarch64.txt

# Link-Time Alignment

```
$ ld.gold -L/usr/lib/gcc-cross/aarch64-linux-gnu/5 -Map map_aarch64.txt \  
main.o fizzbuzz.o -o fizzbuzz_aarch64 \  
<install>/aarch64/lib/crt1.o <install>/aarch64/lib/libc.a \  
<install>/aarch64/lib/libmigrate.a <install>/aarch64/lib/libstack-transform.a \  
<install>/aarch64/lib/libelf.a <install>/aarch64/lib/libc.a \  
--start-group -lgcc -lgcc_eh -end-goup
```



Archive member included because of file (symbol)

map\_aarch64.txt

...

Memory map

```
** file header  
0x0000000000040000 0x40  
** segment headers  
0x0000000000040040 0xe0  
  
.text 0x00000000000400120 0x21744  
.text.exit 0x00000000000400120 0x24 0x4 /home/rlyerly/Downloads/popcorn-compiler/test-  
install/aarch64/lib/libc.a(exit.o) 0x00000000000400120 exit  
.text 0x00000000000400144 0x0 0x4 main.o  
.text.main 0x00000000000400144 0x7c 0x4 main.o  
0x00000000000400144 main  
...
```

# Link-Time Alignment

```
$ ld.gold -L/usr/lib/gcc-cross/aarch64-linux-gnu/5 -Map map_aarch64.txt \  
main.o fizzbuzz.o -o fizzbuzz_aarch64 \  
<install>/aarch64/lib/crt1.o <install>/aarch64/lib/libc.a \  
<install>/aarch64/lib/libmigrate.a <install>/aarch64/lib/libstack-transform.a \  
<install>/aarch64/lib/libelf.a <install>/aarch64/lib/libc.a \  
--start-group -lgcc -lgcc_eh -end-goup
```



Archive member included because of file (symbol)

...  
Memory map

map\_aarch64.txt

```
** File header  
Section name  
00 0x40  
0x0000000000000040 0xe0  
...  
.text 0x000000000000400120 0x21744  
.text.exit 0x000000000000400120 0x24 0x4 /home/rlyerly/Downloads/popcorn-compiler/test-  
install/aarch64/lib/libc.a(exit.o) exit  
0x000000000000400120  
.text 0x000000000000400144 0x0 0x4 main.o  
.text.main 0x000000000000400144 0x7c 0x4 main.o  
0x000000000000400144 main  
...
```

# Link-Time Alignment

```
$ ld.gold -L/usr/lib/gcc-cross/aarch64-linux-gnu/5 -Map map_aarch64.txt \  
main.o fizzbuzz.o -o fizzbuzz_aarch64 \  
<install>/aarch64/lib/crt1.o <install>/aarch64/lib/libc.a \  
<install>/aarch64/lib/libmigrate.a <install>/aarch64/lib/libstack-transform.a \  
<install>/aarch64/lib/libelf.a <install>/aarch64/lib/libc.a \  
--start-group -lgcc -lgcc_eh -end-goup
```



Archive member included because of file (symbol)

...

Memory map

\*\* file header

0x0

\*\* segment headers

0x00000000

0x40

0x80

Virtual memory address

.text	0x0000000000400120	0x21744	
.text.exit	0x0000000000400120	0x24 0x4	/home/rlyerly/Downloads/popcorn-compiler/test-
install/aarch64/lib/libc.a(exit.o)	0x0000000000400120		exit
.text	0x0000000000400144	0x0 0x4	main.o
.text.main	0x0000000000400144	0x7c 0x4	main.o
	0x0000000000400144		main

...

map\_aarch64.txt

# Link-Time Alignment

```
$ ld.gold -L/usr/lib/gcc-cross/aarch64-linux-gnu/5 -Map map_aarch64.txt \  
main.o fizzbuzz.o -o fizzbuzz_aarch64 \  
<install>/aarch64/lib/crt1.o <install>/aarch64/lib/libc.a \  
<install>/aarch64/lib/libmigrate.a <install>/aarch64/lib/libstack-transform.a \  
<install>/aarch64/lib/libelf.a <install>/aarch64/lib/libc.a \  
--start-group -lgcc -lgcc_eh -end-goup
```



Archive member included because of file (symbol)

map\_aarch64.txt

...  
Memory map

** file header		
	0x0000000000400000	0
** segment headers		
	0x0000000000400040	0x0
.text	0x0000000000400120	0x21744
.text.exit	0x0000000000400120	0x24 0x4 /home/rlyerly/Downloads/popcorn-compiler/test-
install/aarch64/lib/libc.a(exit.o)	0x0000000000400120	exit
.text	0x0000000000400144	0x0 0x4 main.o
.text.main	0x0000000000400144	0x7c 0x4 main.o
	0x0000000000400144	main

Size

...

# Link-Time Alignment

```
$ ld.gold -L/usr/lib/gcc-cross/aarch64-linux-gnu/5 -Map map_aarch64.txt \  
main.o fizzbuzz.o -o fizzbuzz_aarch64 \  
<install>/aarch64/lib/crt1.o <install>/aarch64/lib/libc.a \  
<install>/aarch64/lib/libmigrate.a <install>/aarch64/lib/libstack-transform.a \  
<install>/aarch64/lib/libelf.a <install>/aarch64/lib/libc.a \  
--start-group -lgcc -lgcc_eh -end-goup
```



Archive member included because of file (symbol)

map\_aarch64.txt

...  
Memory map

```
** file header  
0x0000000000400000 0x40  
** segment headers  
0x0000000000400040 0xe0  
  
.text 0x0000000000400120 0x21744  
.text.exit 0x0000000000400120 0x24 0x4 /home/rlyerly/Downloads/popcorn-compiler/test-  
install/aarch64/lib/libc.a(exit.o) 0x0000000000400120 exit  
.text 0x0000000000400144 0x0 0x4 main.o  
.text.main 0x0000000000400144 0x7c 0x4 main.o  
0x0000000000400144 main
```

Alignment

...

# Link-Time Alignment

---

```
$ pyalign --compiler-inst <install> \  
  --arm-bin fizzbuzz_aarch64 --arm-map map_aarch64.txt \  
  --x86-bin fizzbuzz_x86_64 --x86-map map_x86_64.txt
```

# Link-Time Alignment

---

```
$ pyalign --compiler-inst <install> \  
  --arm-bin fizzbuzz_aarch64 --arm-map map_aarch64.txt \  
  --x86-bin fizzbuzz_x86_64 --x86-map map_x86_64.txt
```



linker\_script\_arm.x

```
/* Default linker script, for normal executables */  
OUTPUT_FORMAT("elf64-littleaarch64", "elf64-bigaarch64",  
              "elf64-littleaarch64")  
OUTPUT_ARCH(aarch64)  
ENTRY(_start)  
SECTIONS  
{  
...  
  .text      : ALIGN(0x100000)  
  {  
    . = . + 1;  
    . = ALIGN(0x1000);  
    ...  
    . = ALIGN(0x10); /* align for .text.main */  
    "main.o"(.text.main); /* size 0x7c */  
    . = ALIGN(0x10); /* align for .text.fizzbuzz */  
    "fizzbuzz.o"(.text.fizzbuzz); /* size 0xfc */  
    ...  
  }  
}
```



# Link-Time Alignment

```
$ pyalign --compiler-inst <install> \  
  --arm-bin fizzbuzz_aarch64 --arm-map map_aarch64.txt \  
  --x86-bin fizzbuzz_x86_64 --x86-map map_x86_64.txt
```



linker\_script\_arm.x

```
/* Default linker script, for normal executables */  
OUTPUT_FORMAT("elf64-littleaarch64", "elf64-littleaarch64",  
              "elf64-littleaarch64")  
OUTPUT_ARCH(aarch64)  
ENTRY(_start)  
SECTIONS  
{  
  ...  
  .text : ALIGN(0x100000)  
  {  
    . = . + 1;  
    . = ALIGN(0x1000);  
    ...  
    . = ALIGN(0x10); /* align for .text.main */  
    "main.o"(.text.main); /* size 0x7c */  
    . = ALIGN(0x10); /* align for .text.fizzbuzz */  
    "fizzbuzz.o"(.text.fizzbuzz); /* size 0xfc */  
    ...  
  }  
}
```

Re-link each binary with  
generated linker script –  
they're now aligned!

# Post-Processing

---

- Need to reorganize LLVM's stackmap section
  - Not conducive for fast runtime acces
    - Variable sized components within a call site record
    - Stackmap sections from multiple files lumped together
  - Need to uniquify call site IDs across compilation units

# Post-Processing

---

```
gen-stackinfo -f fizzbuzz_aarch64 && readelf -SW fizzbuzz_aarch64
```

# Post-Processing

```
gen-stackinfo -f fizzbuzz_aarch64 && readelf -SW fizzbuzz_aarch64
```



There are 38 section headers, starting at offset 0x66a3cb:

Section Headers:

[Nr]	Name	Type	Address	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	0000000000000000	000000	000000	00		0	0	0
[ 1]	.text	PROGBITS	0000000000500000	100000	024ec4	00	AX	0	0	1048576
[ 2]	.rodata	PROGBITS	0000000000600000	200000	00438f	00	A	0	0	1048576
...										
[31]	.symtab	SYMTAB	0000000000000000	655618	00fe10	18		32	2280	8
[32]	.strtab	STRTAB	0000000000000000	665428	0039ef	00		0	0	1
[33]	.shstrtab	STRTAB	0000000000000000	668e17	000208	00		0	0	1
[34]	.stack_transform.id	PROGBITS	0000000000000000	66901f	0007ec	34		0	0	0
[35]	.stack_transform.addr	PROGBITS	0000000000000000	66980b	0007ec	34		0	0	0
[36]	.stack_transform.live	PROGBITS	0000000000000000	669ff7	000348	0c		0	0	0
[37]	.stack_transform.arch_const	PROGBITS	0000000000000000	66a33f	00008c	14		0	0	0

# Post-Processing

```
gen-stackinfo -f fizzbuzz_aarch64 && readelf -SW fizzbuzz_aarch64
```



There are 38 section headers, starting at offset 0x66a3cb:

Section Headers:

[Nr]	Name	Type	Address	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	0000000000000000	000000	000000	00		0	0	0
[ 1]	.text	PROGBITS	0000000000500000	100000	024ec4	00	AX	0	0	1048576
[ 2]	.rodata							0	0	1048576
...										
[31]	.symtab							32	2280	8
[32]	.strtab	STRTAB	0000000000000000	665428	0039ef	00		0	0	1
[33]	.shstrtab	SHSTRTAB	0000000000000000	668e17	000208	00		0	0	1
[34]	.stack_transform.id	PROGBITS	0000000000000000	66901f	0007ec	34		0	0	0
[35]	.stack_transform.addr	PROGBITS	0000000000000000	66980b	0007ec	34		0	0	0
[36]	.stack_transform.live	PROGBITS	0000000000000000	669ff7	000348	0c		0	0	0
[37]	.stack_transform.arch_const	PROGBITS	0000000000000000	66a33f	00008c	14		0	0	0

Consumed by the stack transformation runtime

# Post-Processing

---

- The binaries are now ready for runtime migration
  - One binary per target in the format `<name>_<target>`
- Useful utilities:
  - `check-align.py`: consume the symbol tables of two binaries to verify that all symbols begin at aligned virtual addresses
  - `check-stackmaps.py`: sanity check the stackmaps generated by the compiler (same numbers and types of functions, call sites, live values)
    - Must be run post-alignment – matches functions by address

---

# **Part 4:**

# **Migration & Stack Transformation**

# Migrating Between Architectures

---

- Migration in Popcorn Linux
  - Somebody (either inside or outside the application) **proposes** that a given thread migrates to a given destination node
  - The `check_migrate` function queries whether a migration has been proposed for the current thread
  - If so, `check_migrate` invokes stack transformation
    - Take a snapshot of current register set
    - Rewrite stack to another location in memory (in userspace)
    - Return populated destination ISA register set
  - Pass rewritten register set to kernel's thread migration service
  - Resume execution in `check_migrate` on destination



# Triggering Migration

---

- Proposing migration (`libmigrate:trigger.c`):

```
syscall(SYS_CALL_PROPOSE_MIGRATION, pid, nid)
```

# Triggering Migration

---

- Proposing migration (`libmigrate:trigger.c`):

Direct shortcut to the system call interface

`syscall(SYS CALL_PROPOSE_MIGRATION, pid, nid)`

# Triggering Migration

---

- Proposing migration (`libmigrate:trigger.c`):

System call number – proposing a migration

```
syscall(SYS CALL_PROPOSE_MIGRATION, pid, nid)
```

# Triggering Migration

---

- Proposing migration (`libmigrate:trigger.c`):

For which task are we proposing migration

```
syscall(SYS CALL_PROPOSE_MIGRATION, pid, nid)
```

# Triggering Migration

---

- Proposing migration (`libmigrate:trigger.c`):

To which node the task should migrate

```
syscall(SYS CALL_PROPOSE_MIGRATION, pid, nid)
```

# Triggering Migration

---

- Checking to see if a migration was proposed (`libmigrate:migrate.c`):

```
syscall(SYS_CALL_MIGRATION_PROPOSED)
```

# Triggering Migration

---

- Checking to see if a migration was proposed (`libmigrate:migrate.c`):

Return whether a migration was proposed for the current thread

- $\geq 0$ : proposed destination node
- $< 0$ : no migration proposed

`syscall(SYSCALL_MIGRATION_PROPOSED)`

# Triggering Migration

---

- Invoking thread migration service  
(`libmigrate:migrate.c`):

```
syscall(SYSCALL_SCHED_MIGRATE, nid, regs_dst, sp, bp)
```



# Triggering Migration

---

- Invoking thread migration service  
(`libmigrate:migrate.c`):

Migrate the thing!

```
syscall(SYSCALL_SCHED_MIGRATE, nid, regs_dst, sp, bp)
```

# Triggering Migration

---

- Invoking thread migration service (`libmigrate:migrate.c`):

Where to migrate

```
syscall(SYSCALL_SCHED_MIGRATE, nid, regs_dst, sp, bp)
```

# Triggering Migration

---

- Invoking thread migration service  
(`libmigrate:migrate.c`):

Destination register set – thread will be restarted with these registers on destination

- Generated by stack transformation runtime

```
syscall(SYSCALL_SCHED_MIGRATE, nid, regs_dst, sp, bp)
```

# Triggering Migration

---

- Invoking thread migration service  
(`libmigrate:migrate.c`):

New stack & frame pointer on destination

```
syscall(SYS_CALL_SCHED_MIGRATE, nid, regs_dst, sp, bp)
```

# Triggering Migration

---

- Invoking thread migration service  
(`libmigrate:migrate.c`):

`syscall ($`

**These interfaces are subject to change – see  
<repo>/lib/migrate for up-to-date versions!**

ation

`sp, bp)`


# Stack Transformation

---

- Rewrite entire stack from outermost frame inwards
- Stack transformation runtime (the “runtime”) opens metadata sections using `libelf`
  - `.stack_transform.unwind` – per-function callee-saved register locations on the stack for frame unwinding
  - `.stack_transform.unwind_arange` – address ranges for functions in the binary (used for bootstrapping outermost frame)
  - `.stack_transform.id` – call sites (stackmaps) sorted by ID
  - `.stack_transform.addr` – call sites sorted by program location
  - `.stack_transform.live` – live value location records
  - `.stack_transform.arch_const` – architecture-specific live value records

# Stack Transformation

- All transformation is performed in userspace
  - Linux by default gives the main thread 8MB of stack space (musl is modified to do the same for spawned threads)
  - Divide stack into upper/lower half, rewrite from one to the other at migration time




# Stack Transformation

---

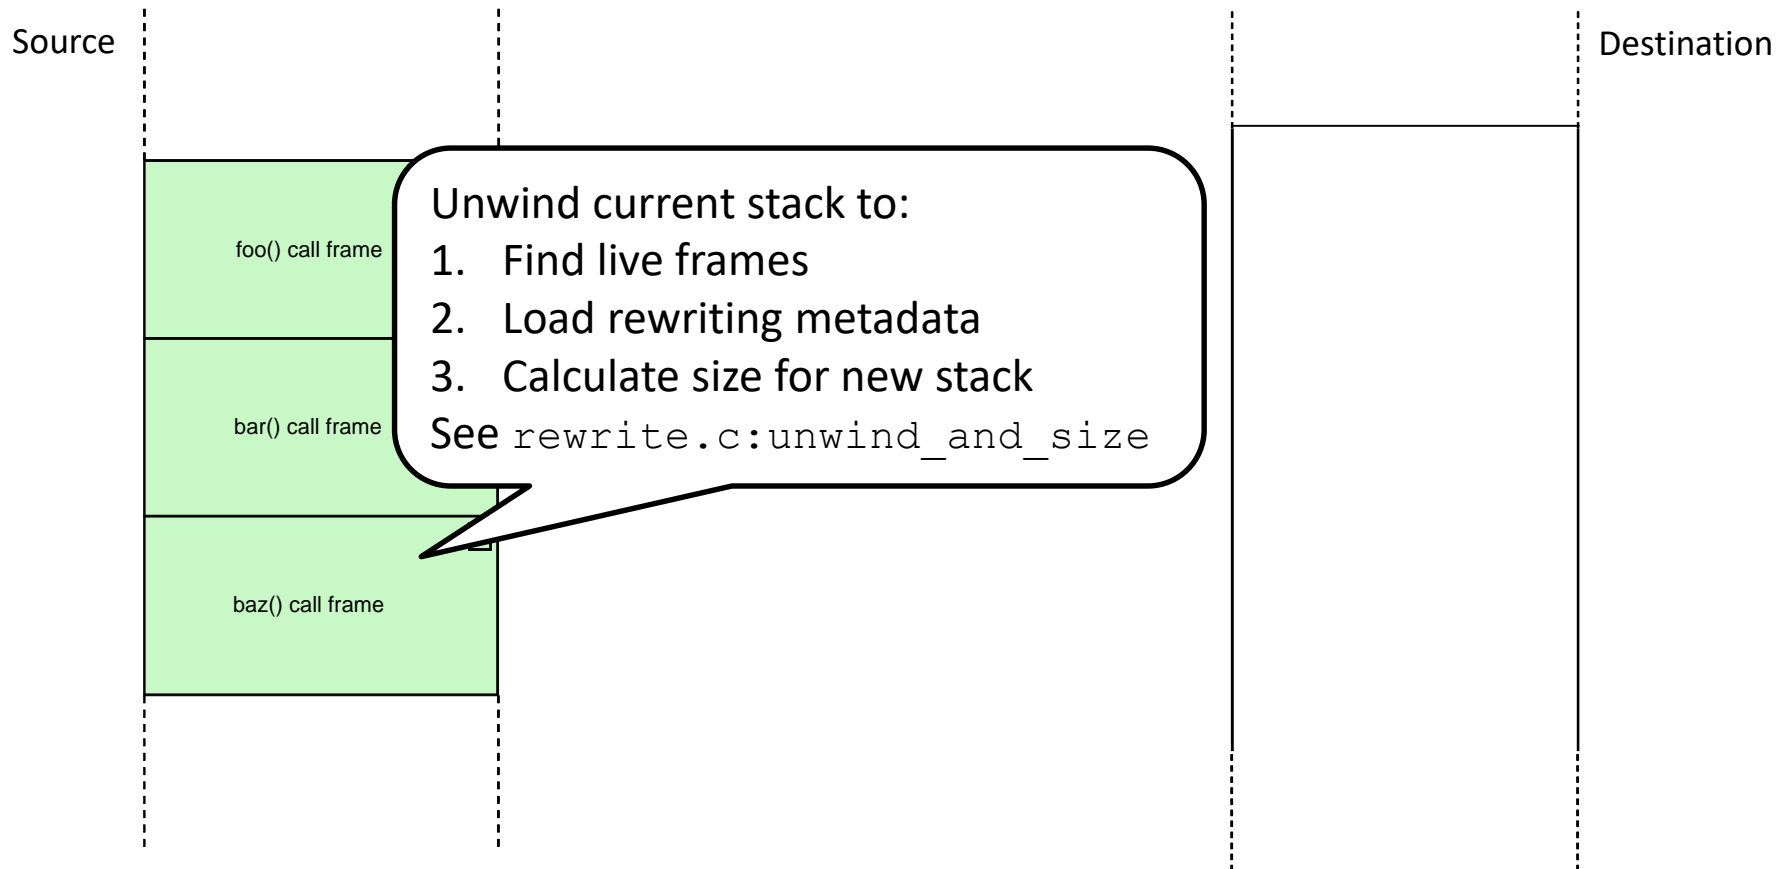
- Runtime initializes handles containing transformation metadata at program startup
  - Ways to tell runtime name of binary containing a target's metadata
    1. Set environment variable `ST_<target>_BIN`
    2. Define symbol `<target>_fn`
    3. If (1) & (2) not used, runtime will default to appending target to current executing binary's name



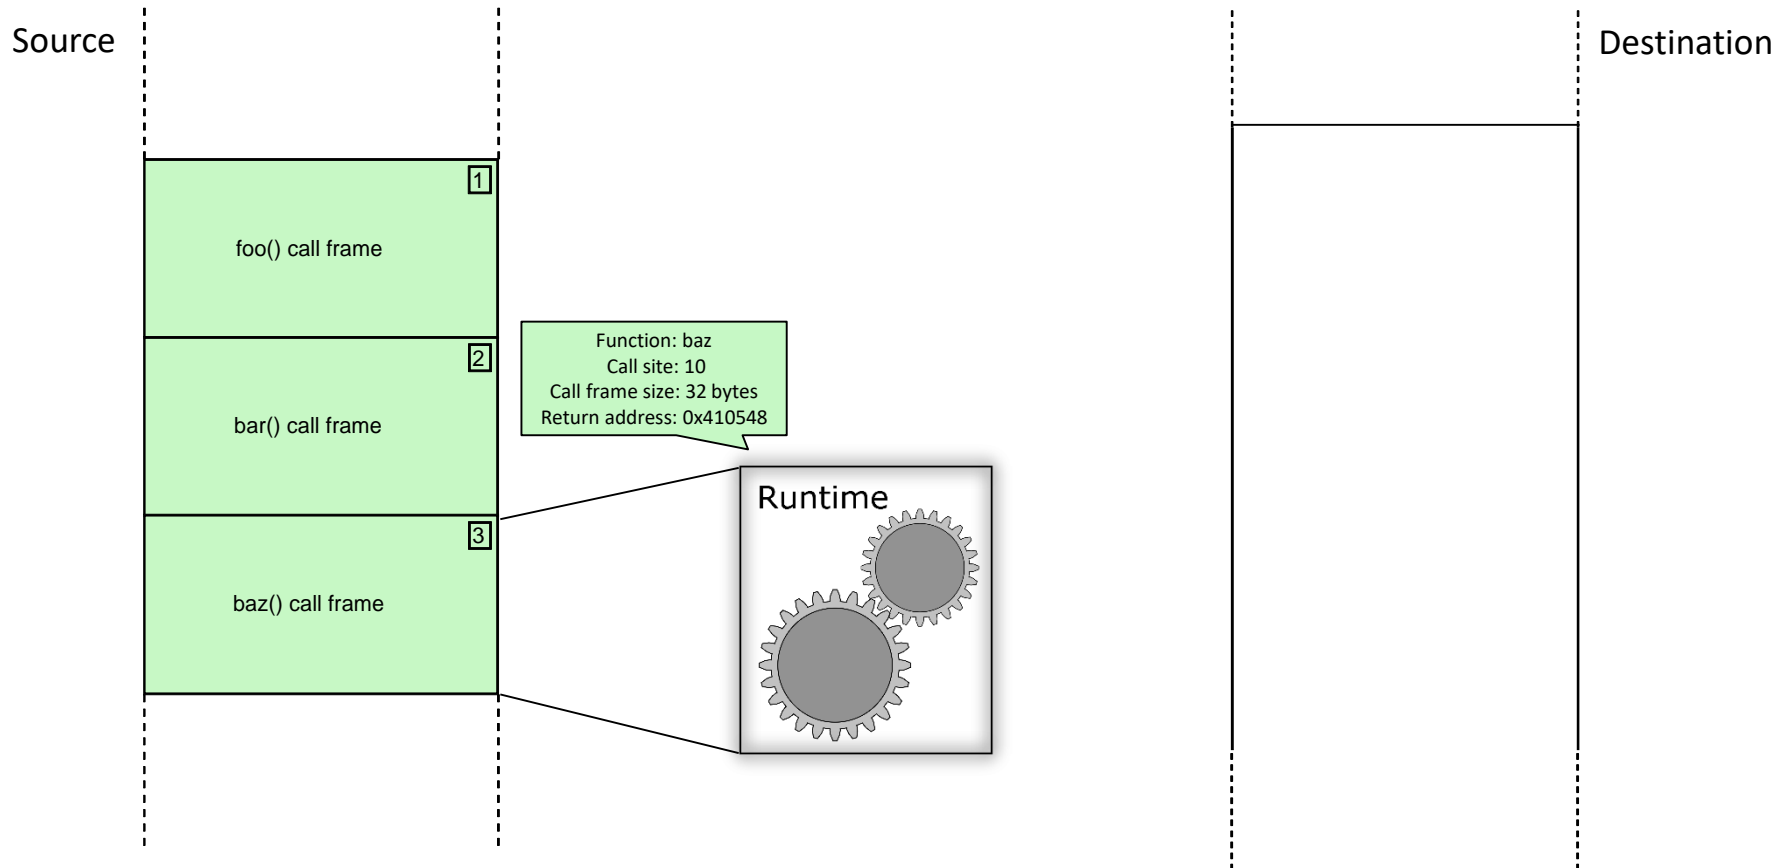
# Stack Transformation



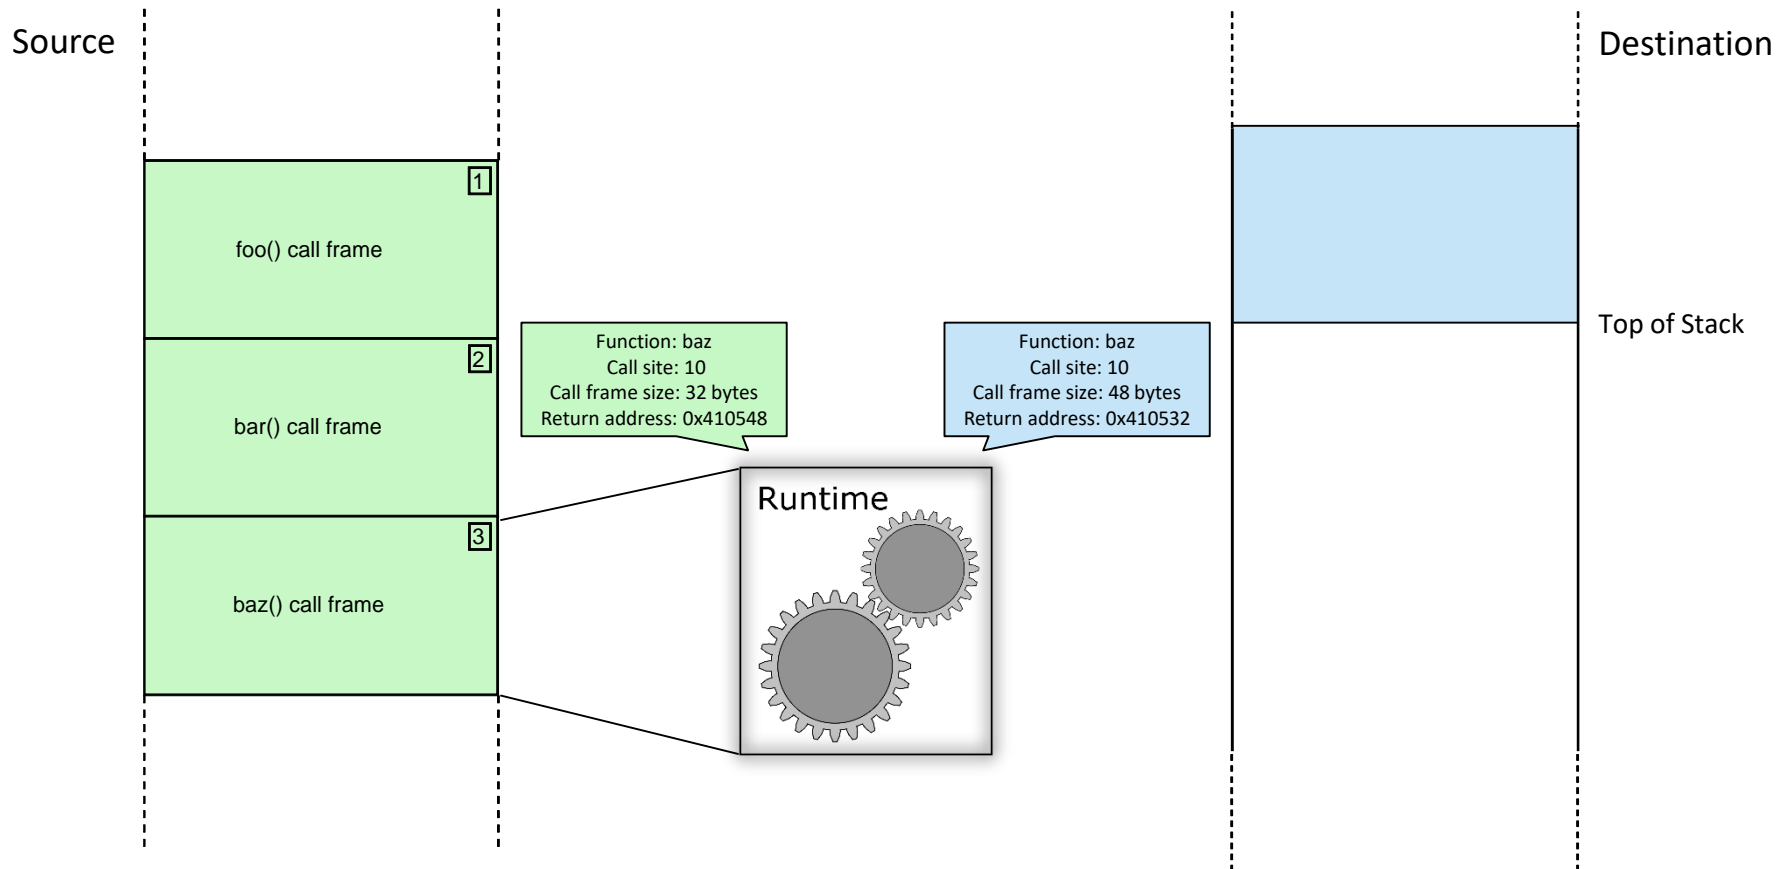
# Stack Transformation



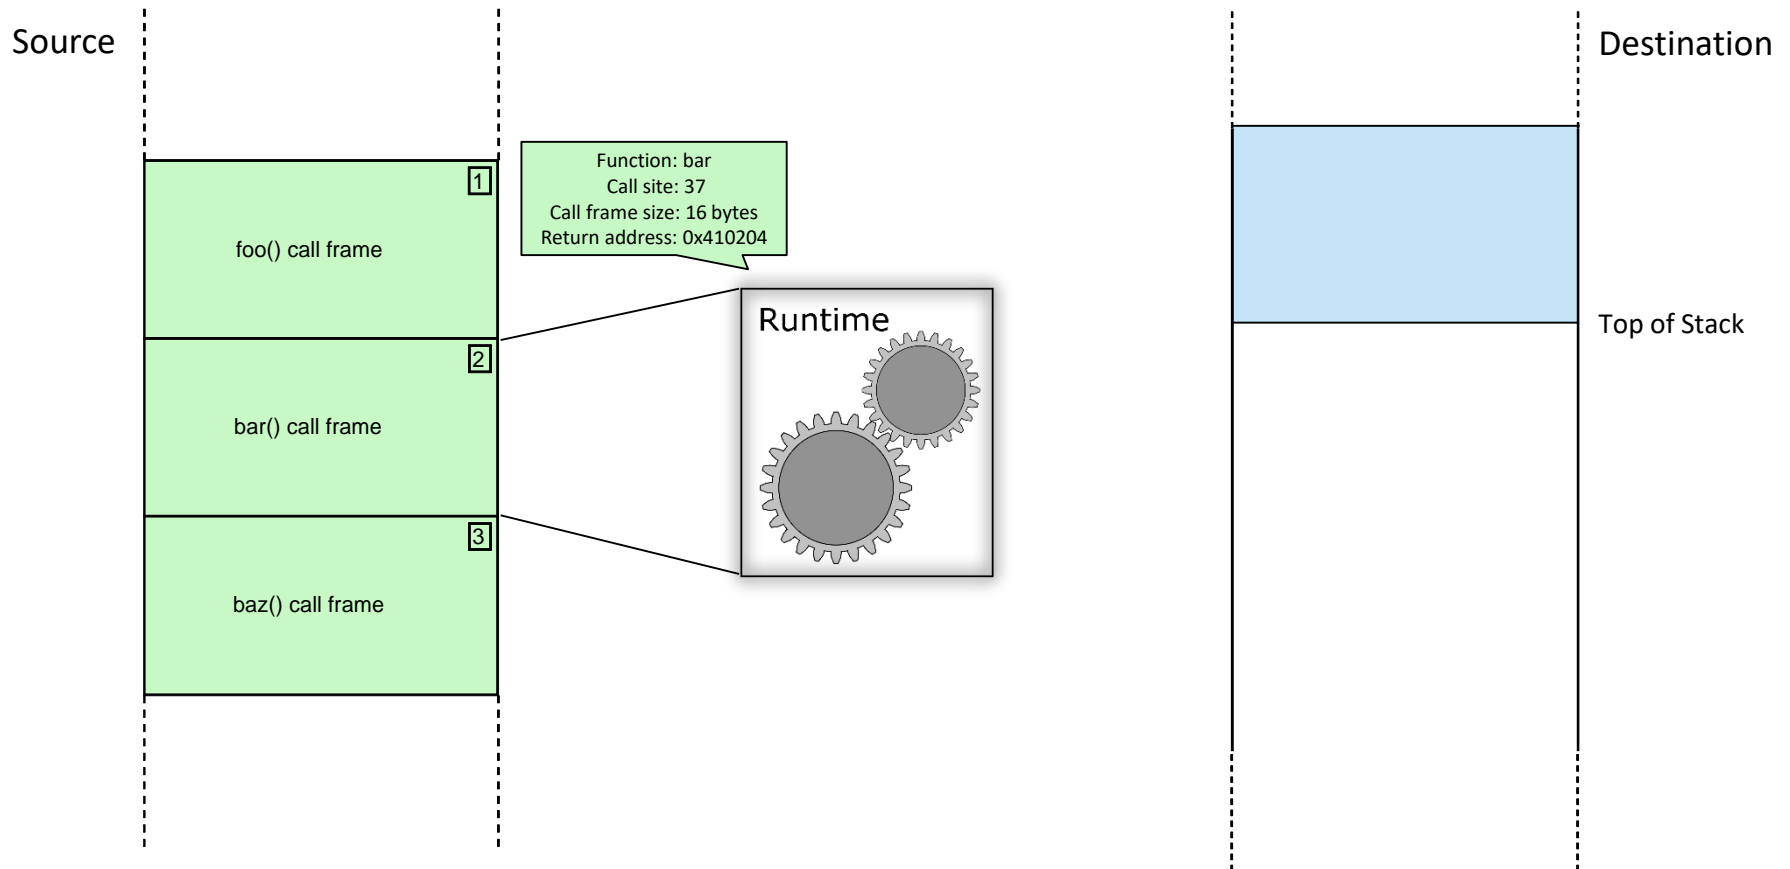
# Stack Transformation



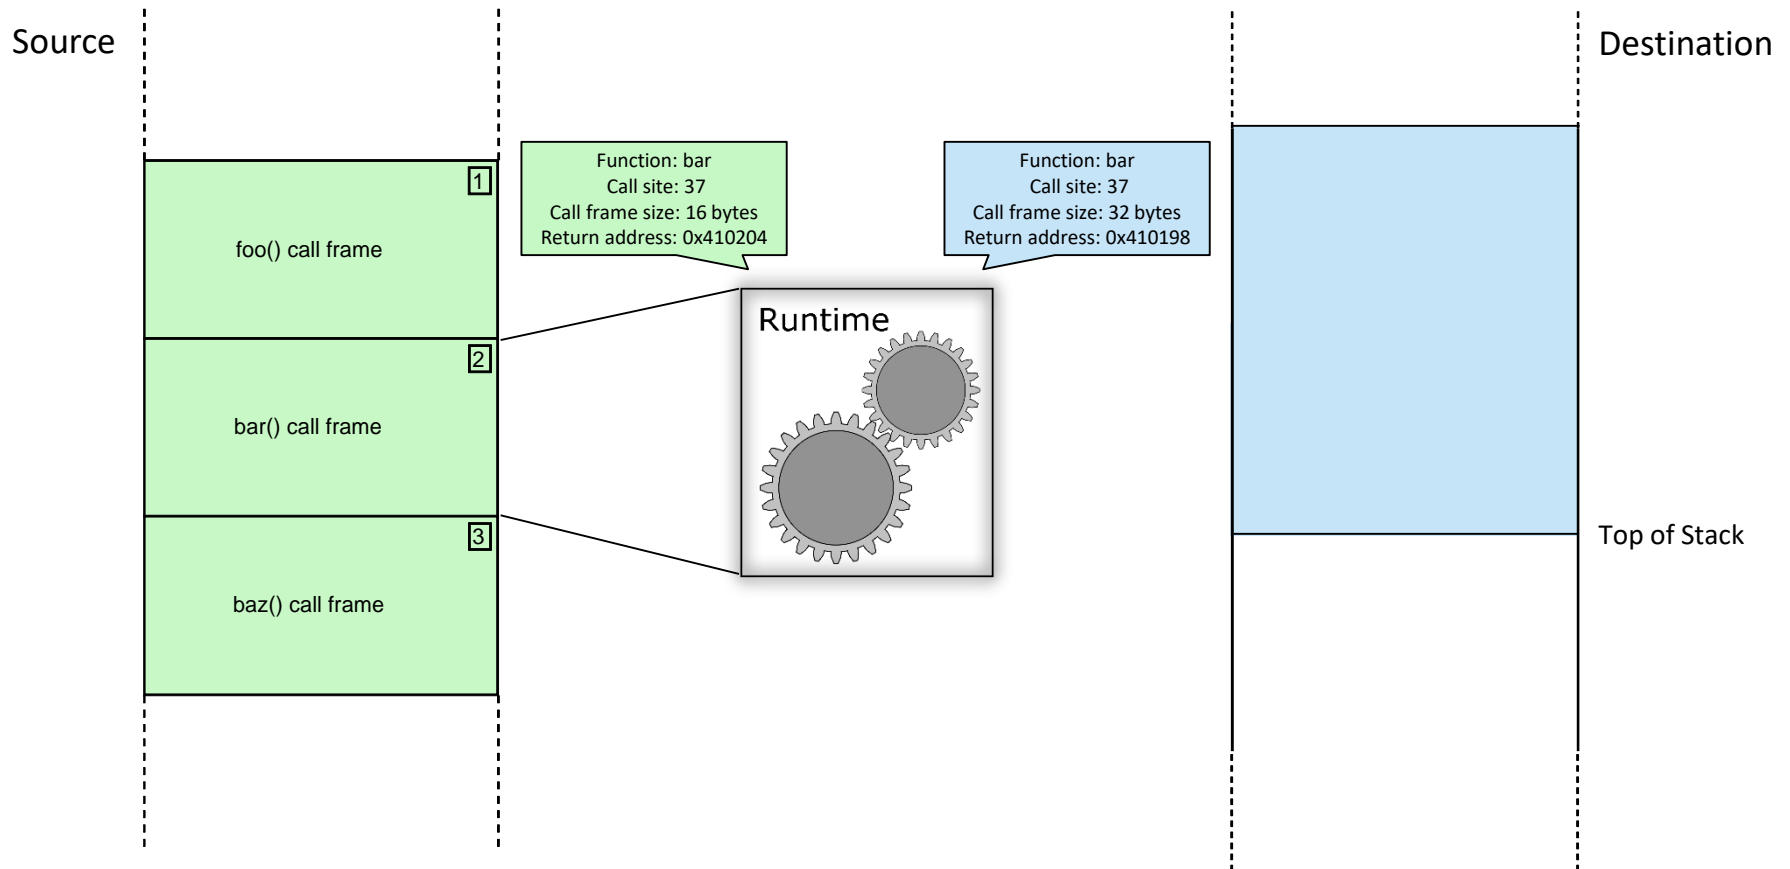
# Stack Transformation



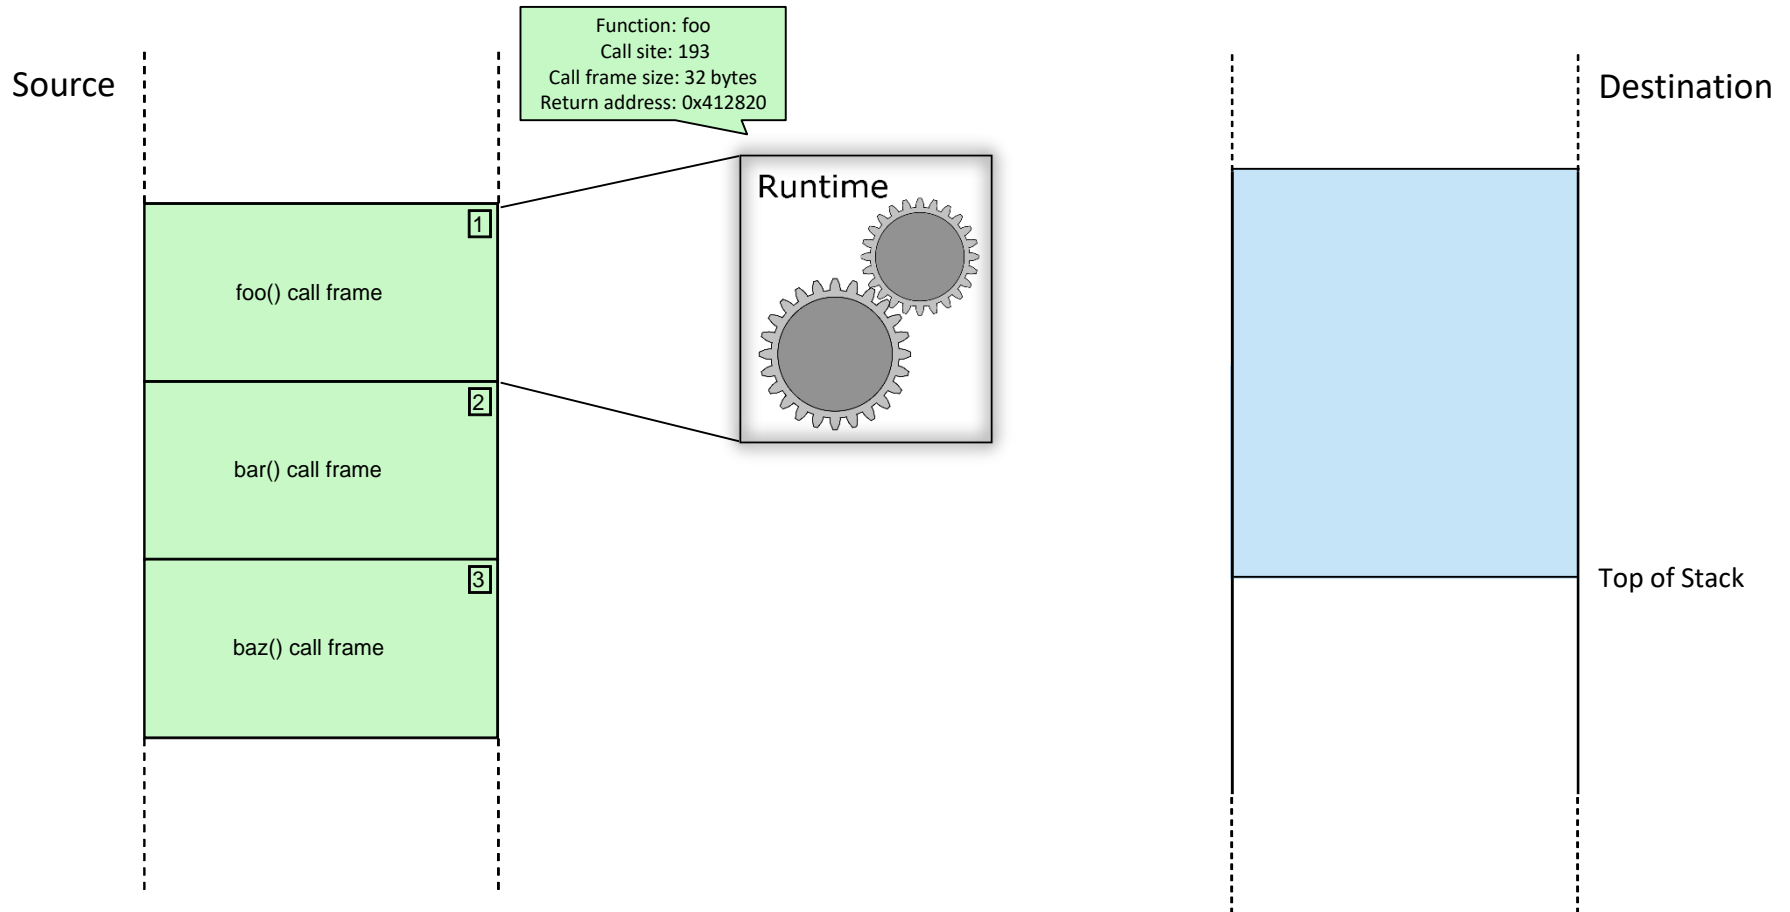
# Stack Transformation



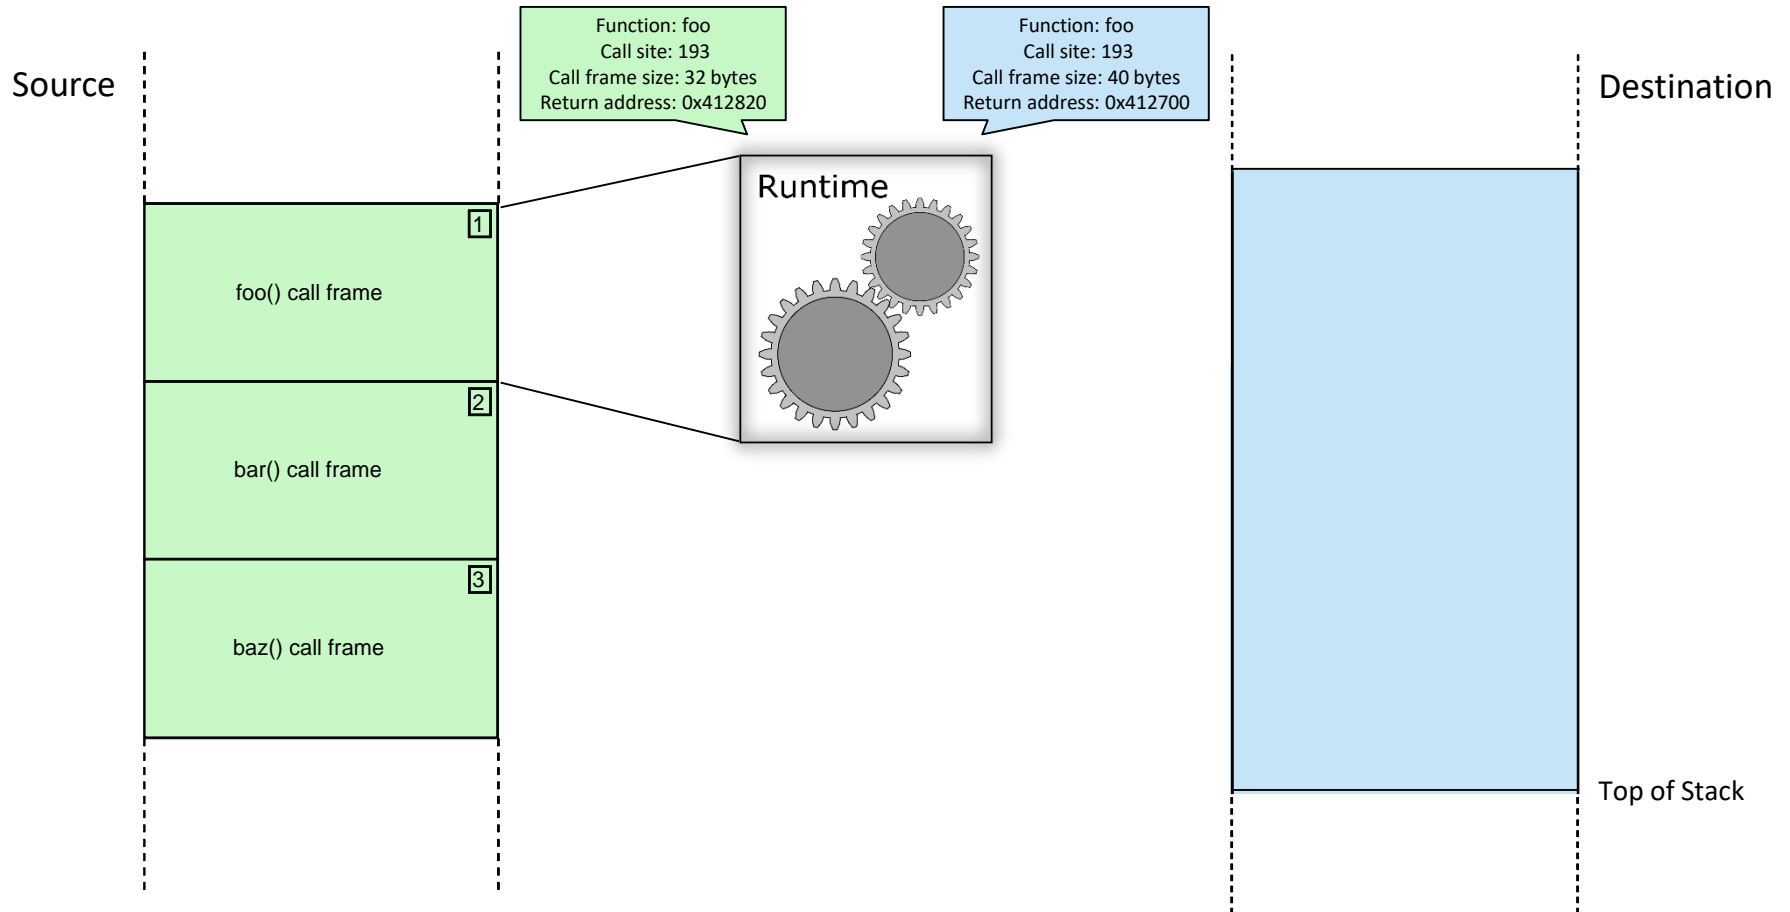
# Stack Transformation



# Stack Transformation

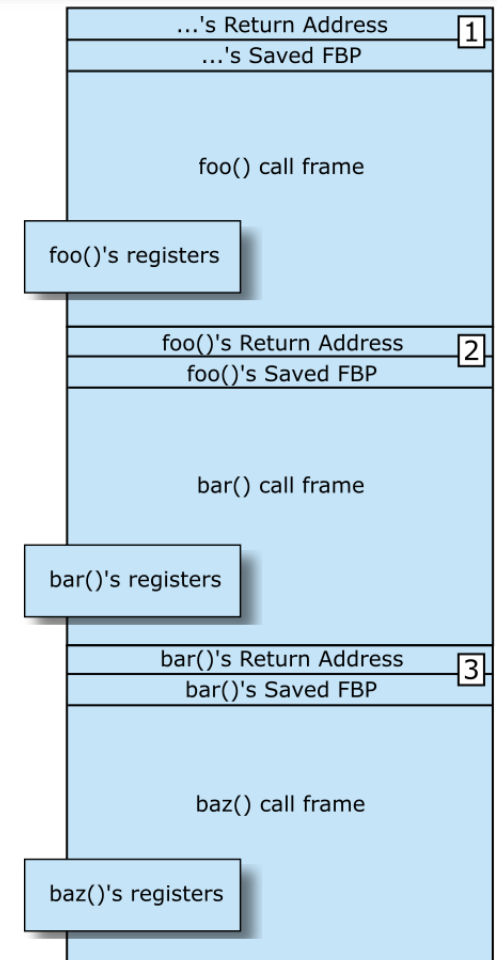

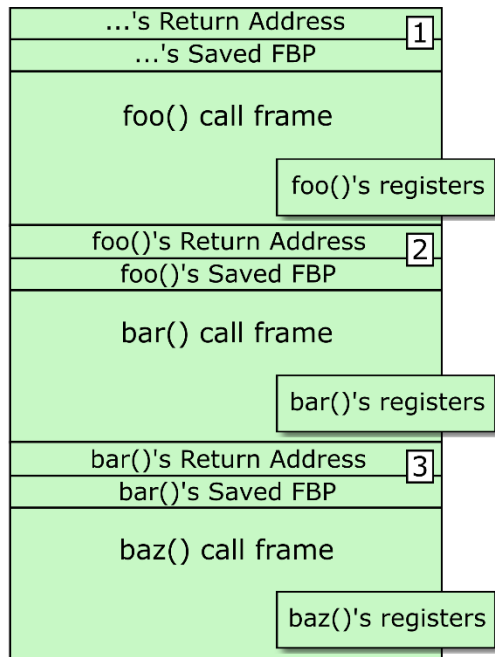


# Stack Transformation


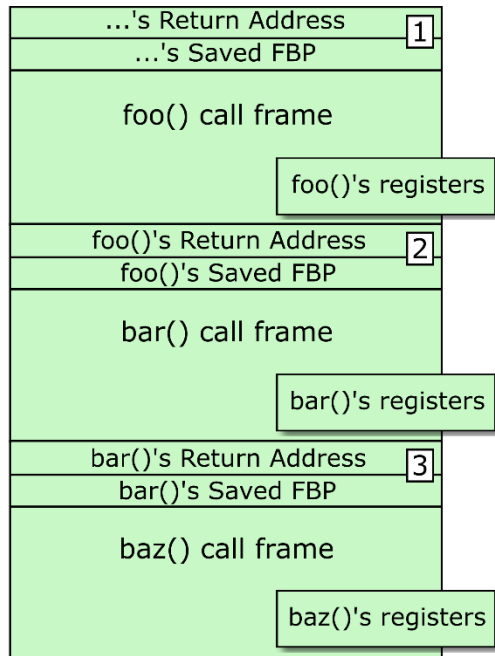




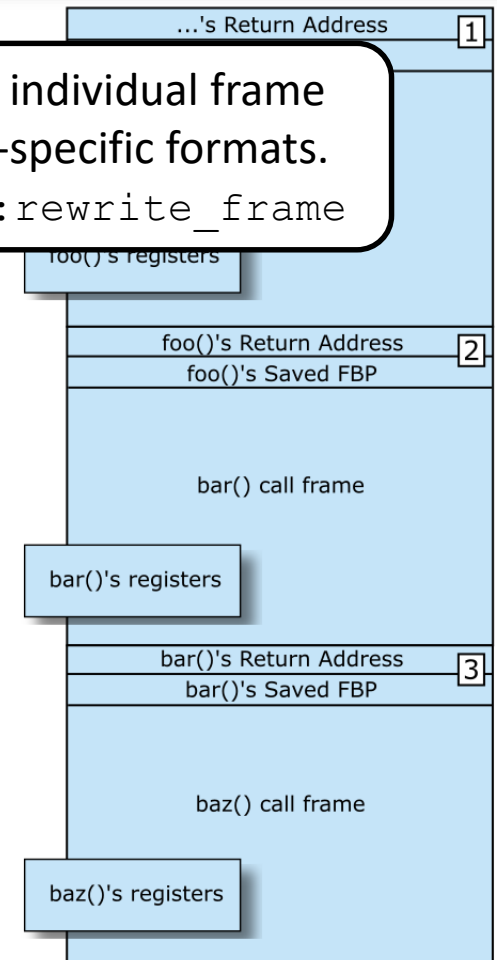
# Stack Transformation



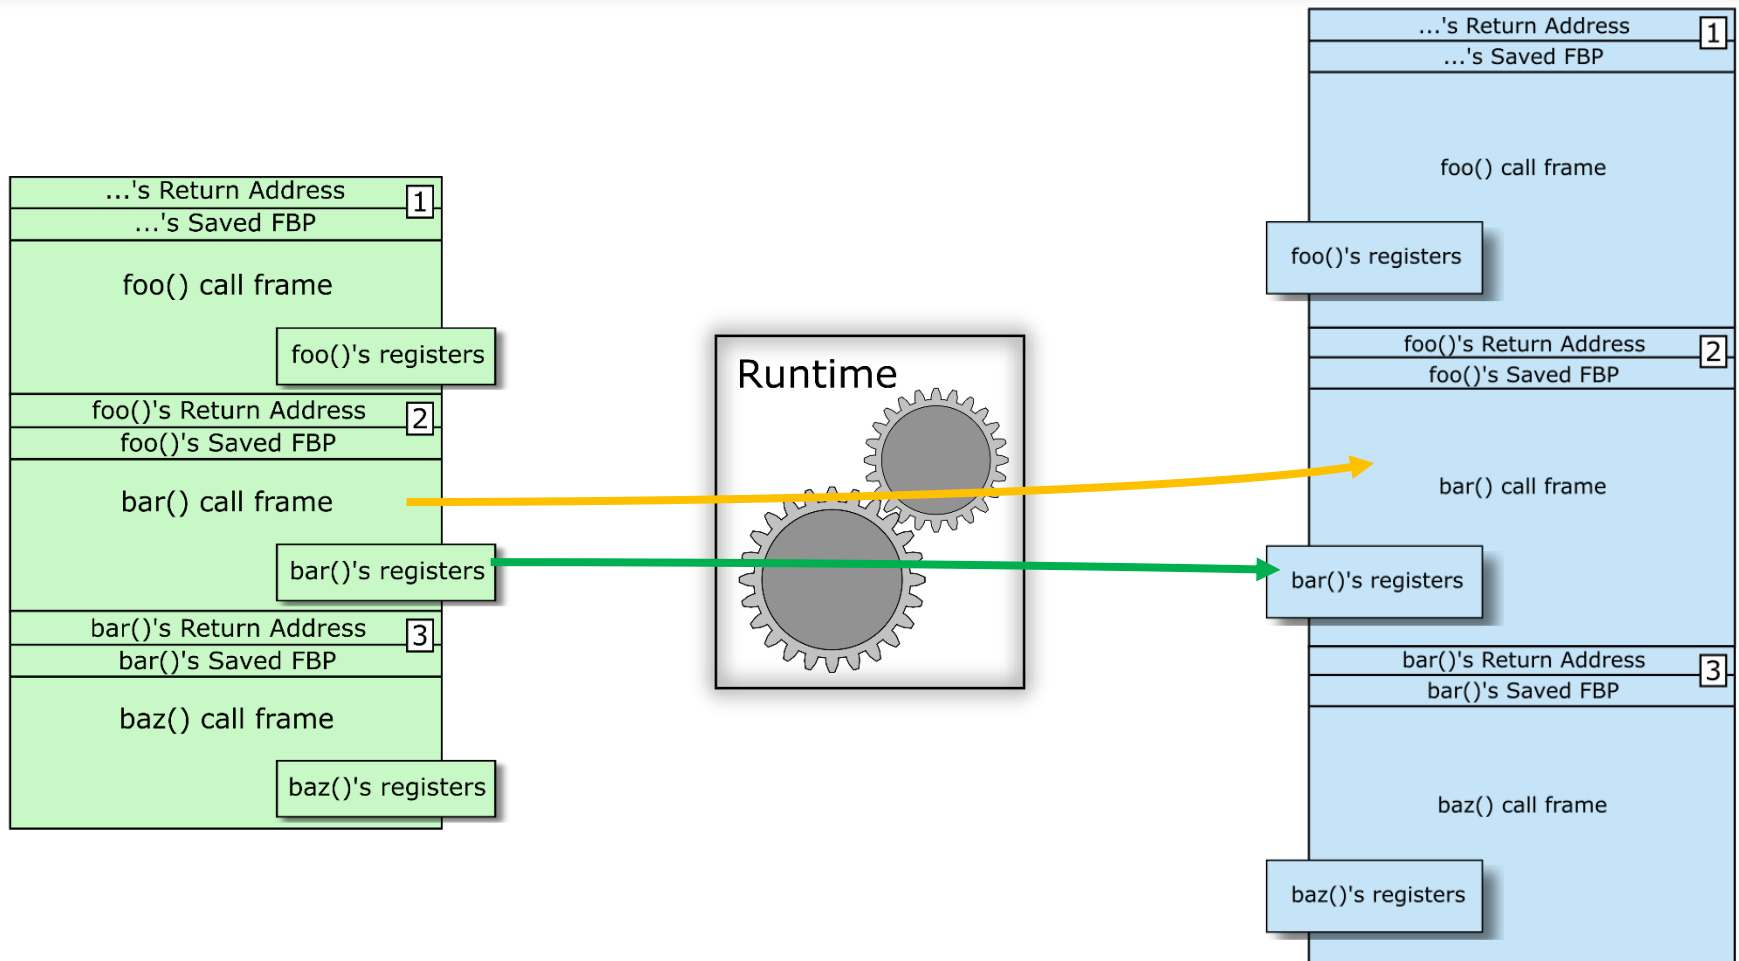
# Stack Transformation





Transform each individual frame between target-specific formats.  
See `rewrite.c:rewrite_frame`



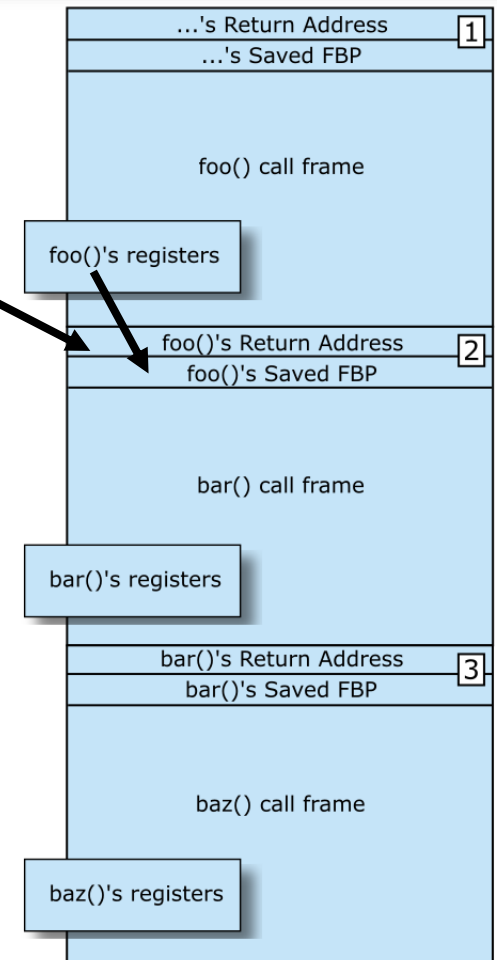
# Stack Transformation





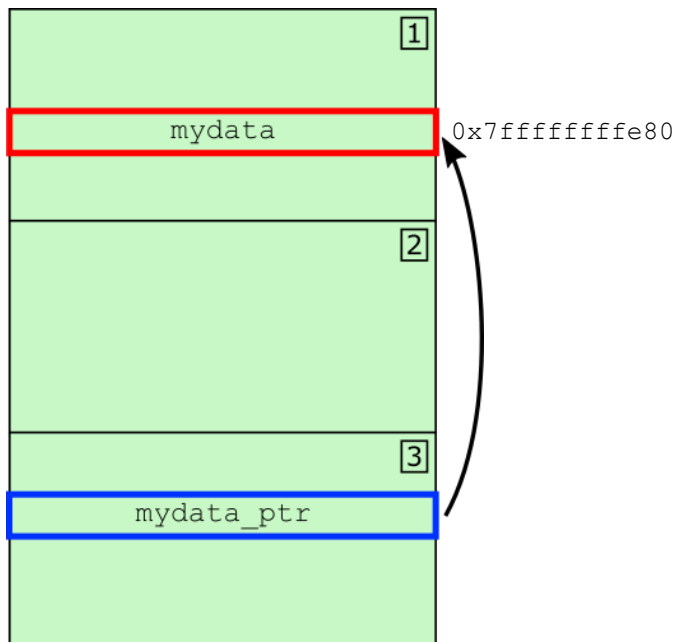
# Stack Transformation



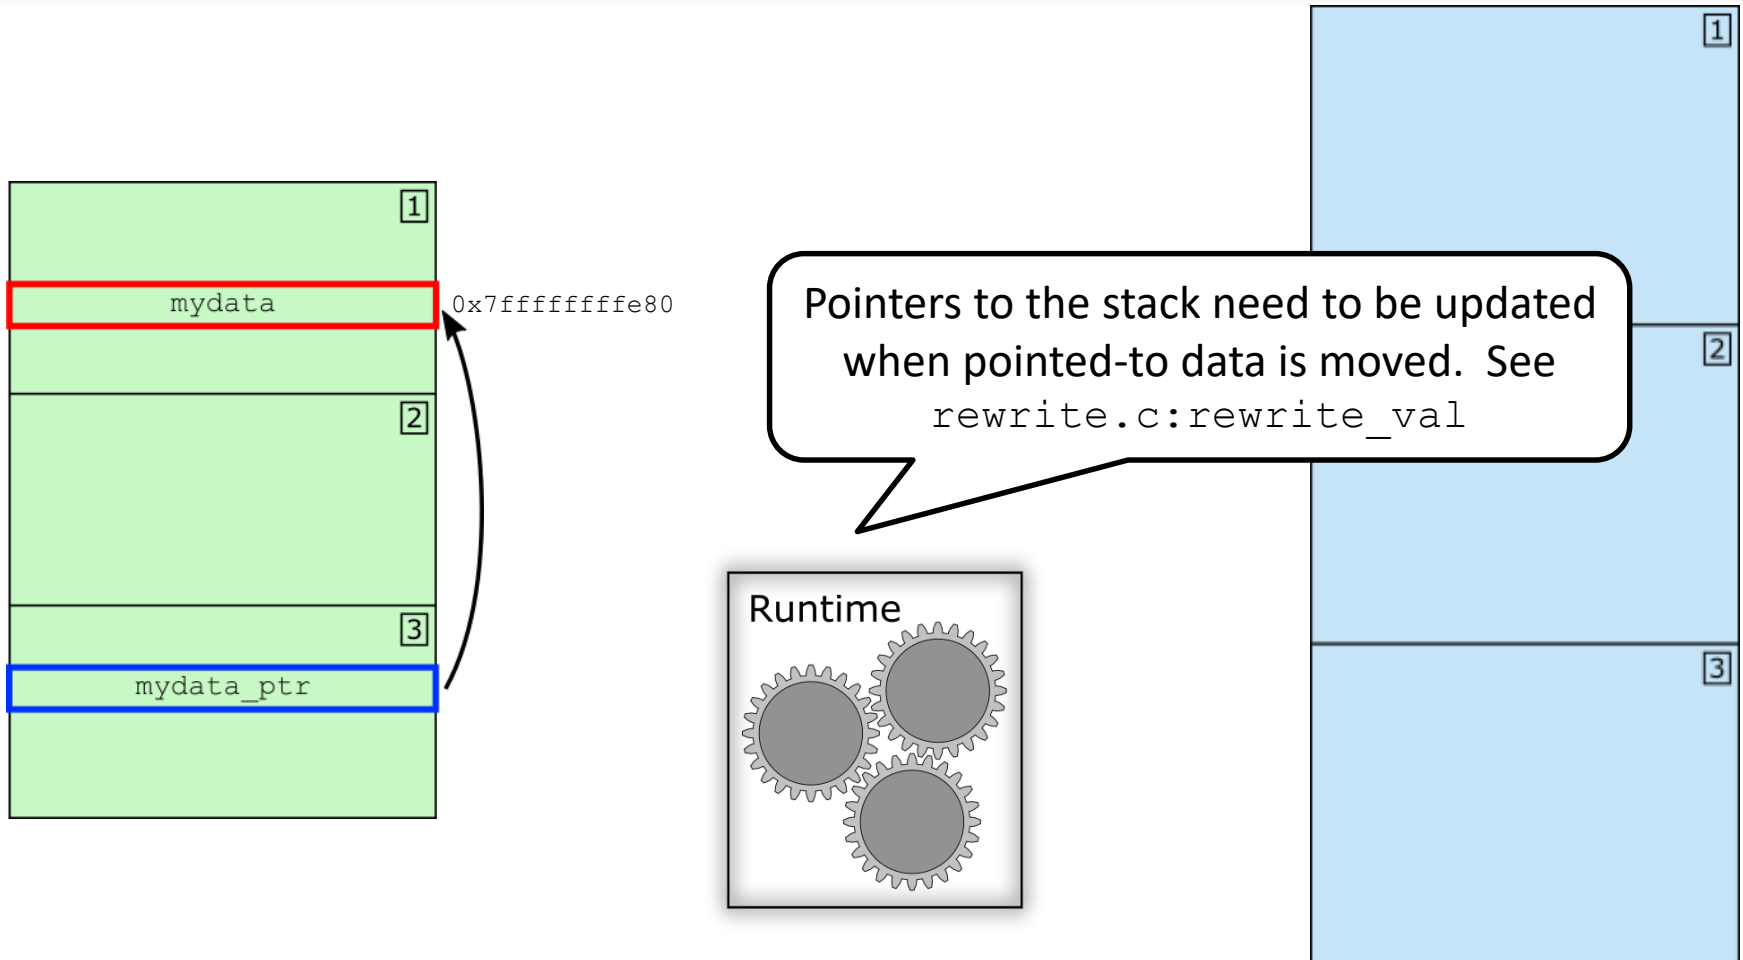
Function: foo  
Call site: 193  
Call frame size: 40 bytes  
Return address: 0x412700



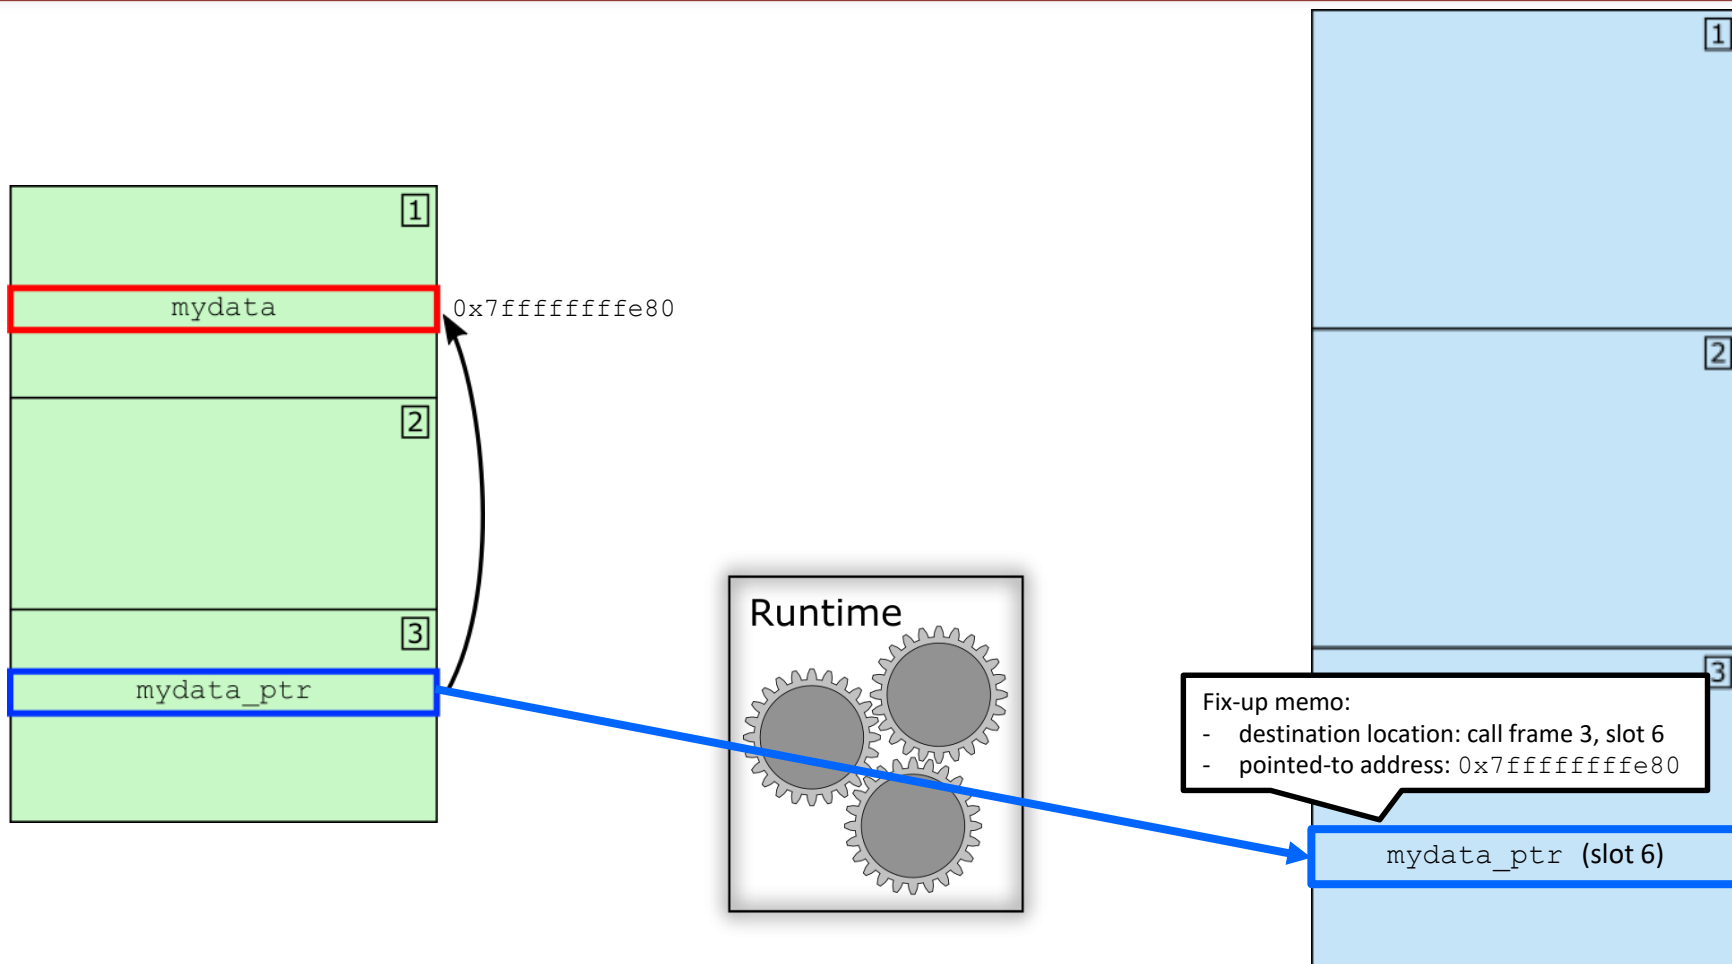
# State Transformation



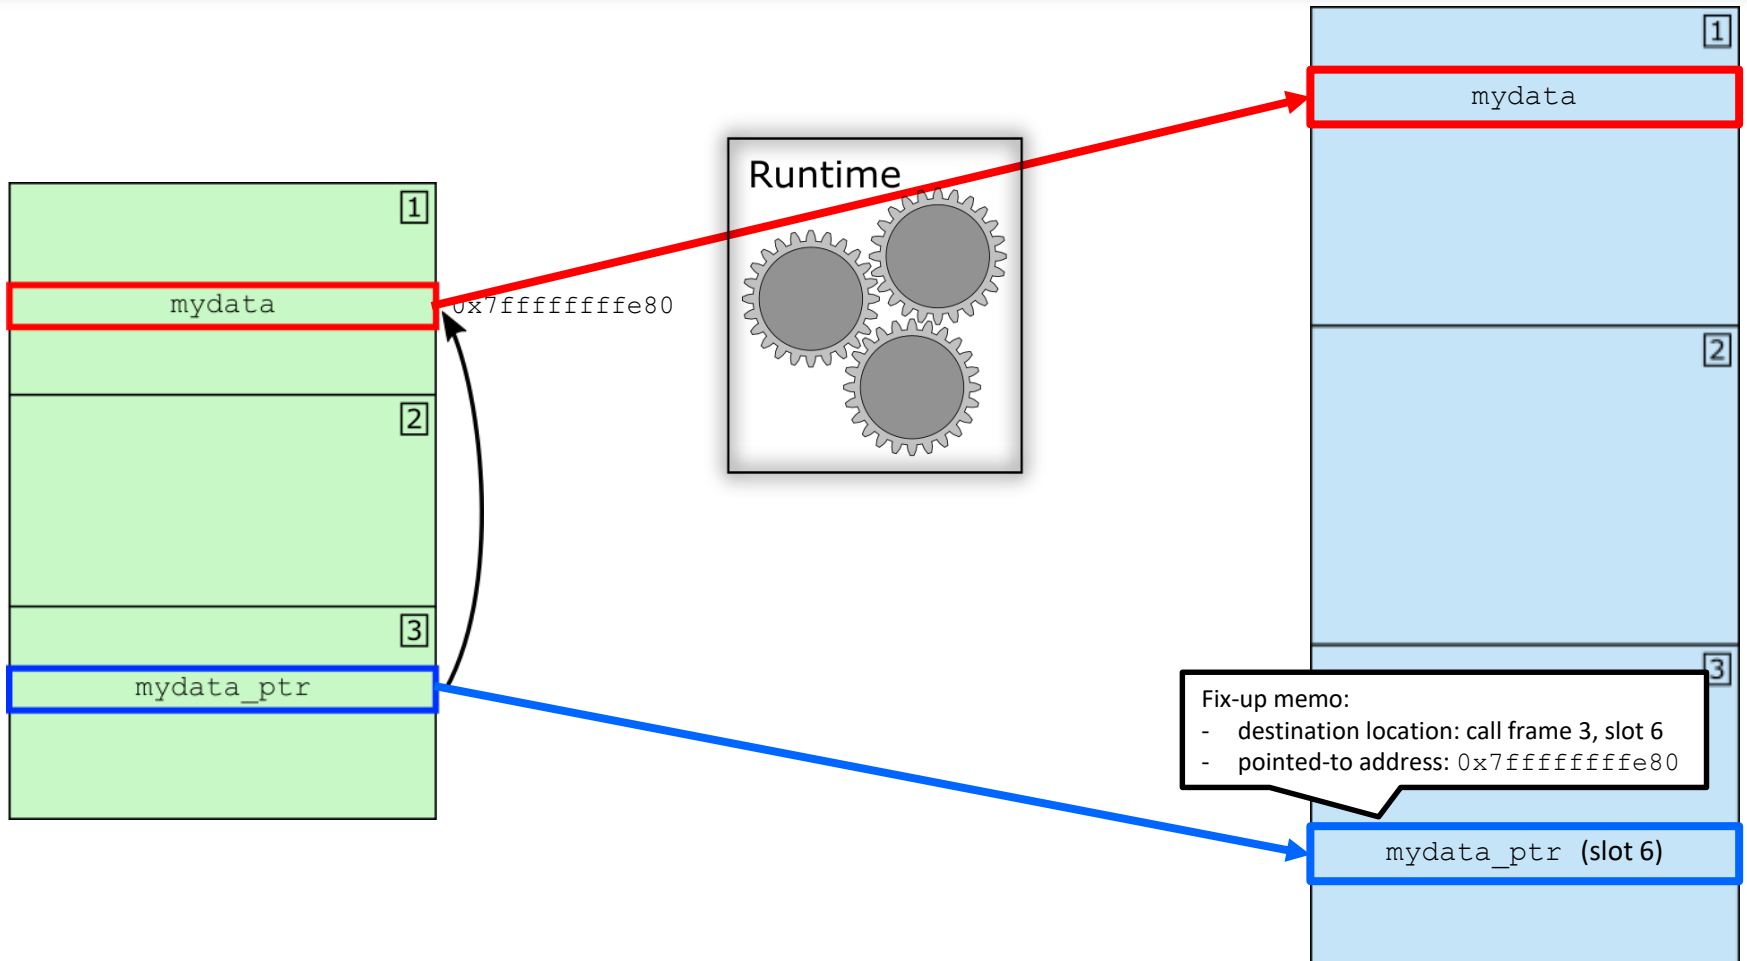
# State Transformation



# State Transformation

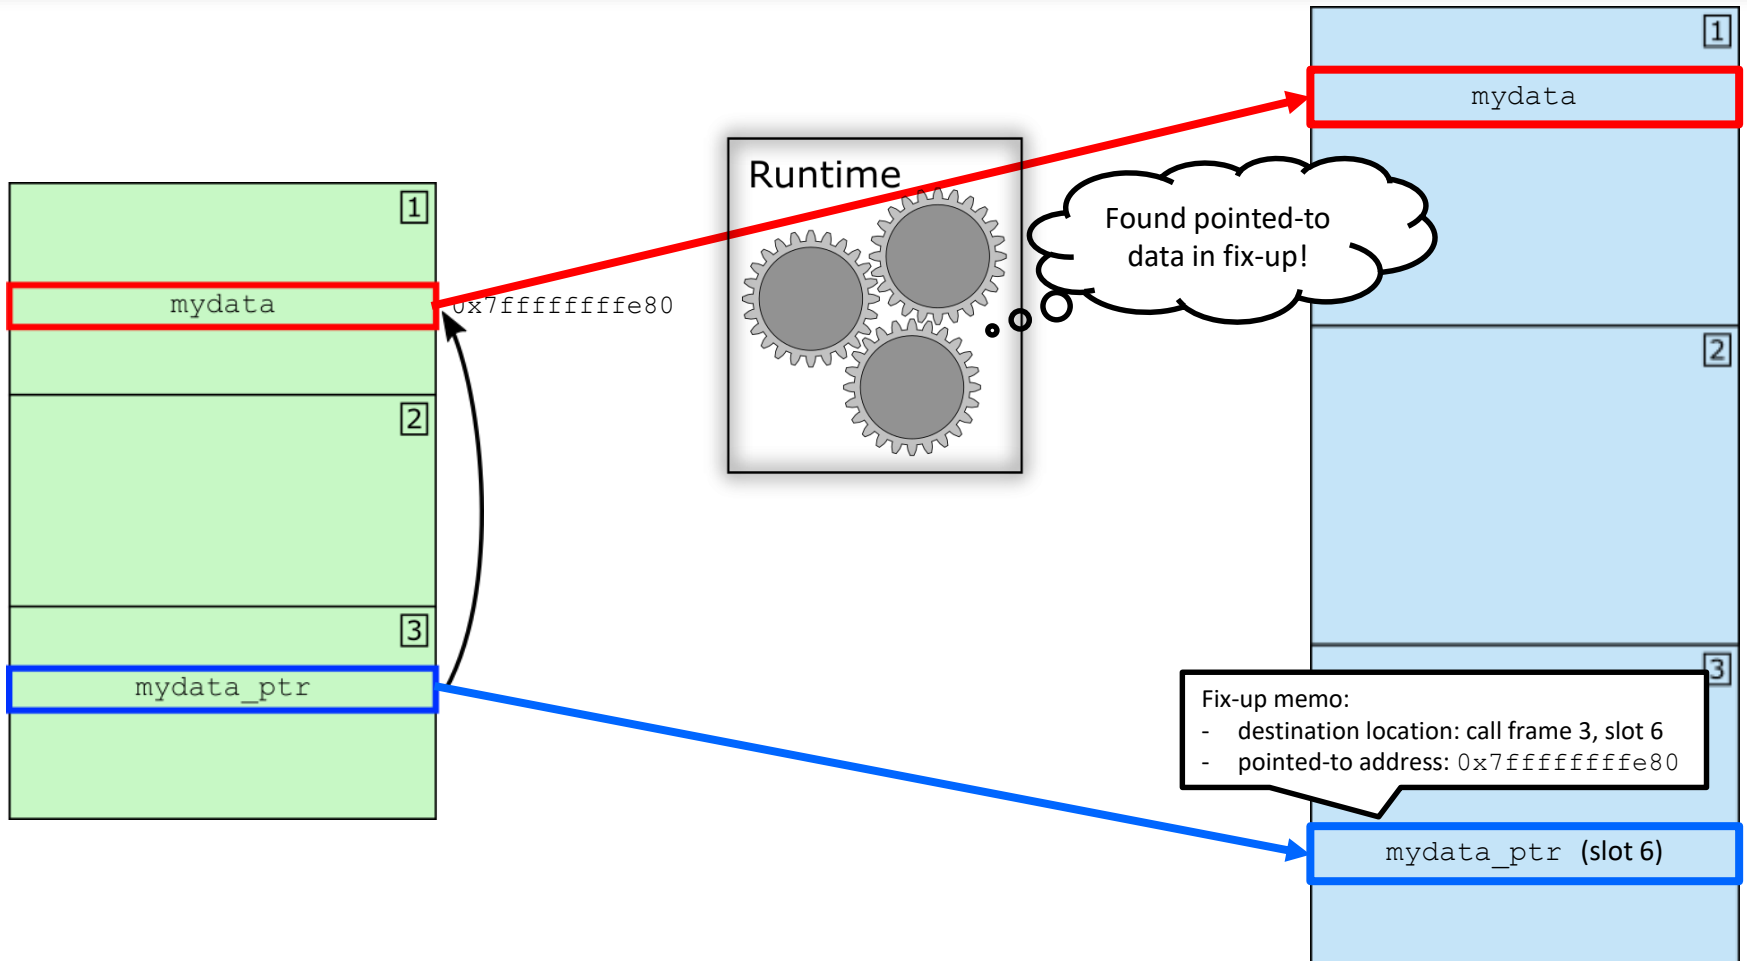


# State Transformation

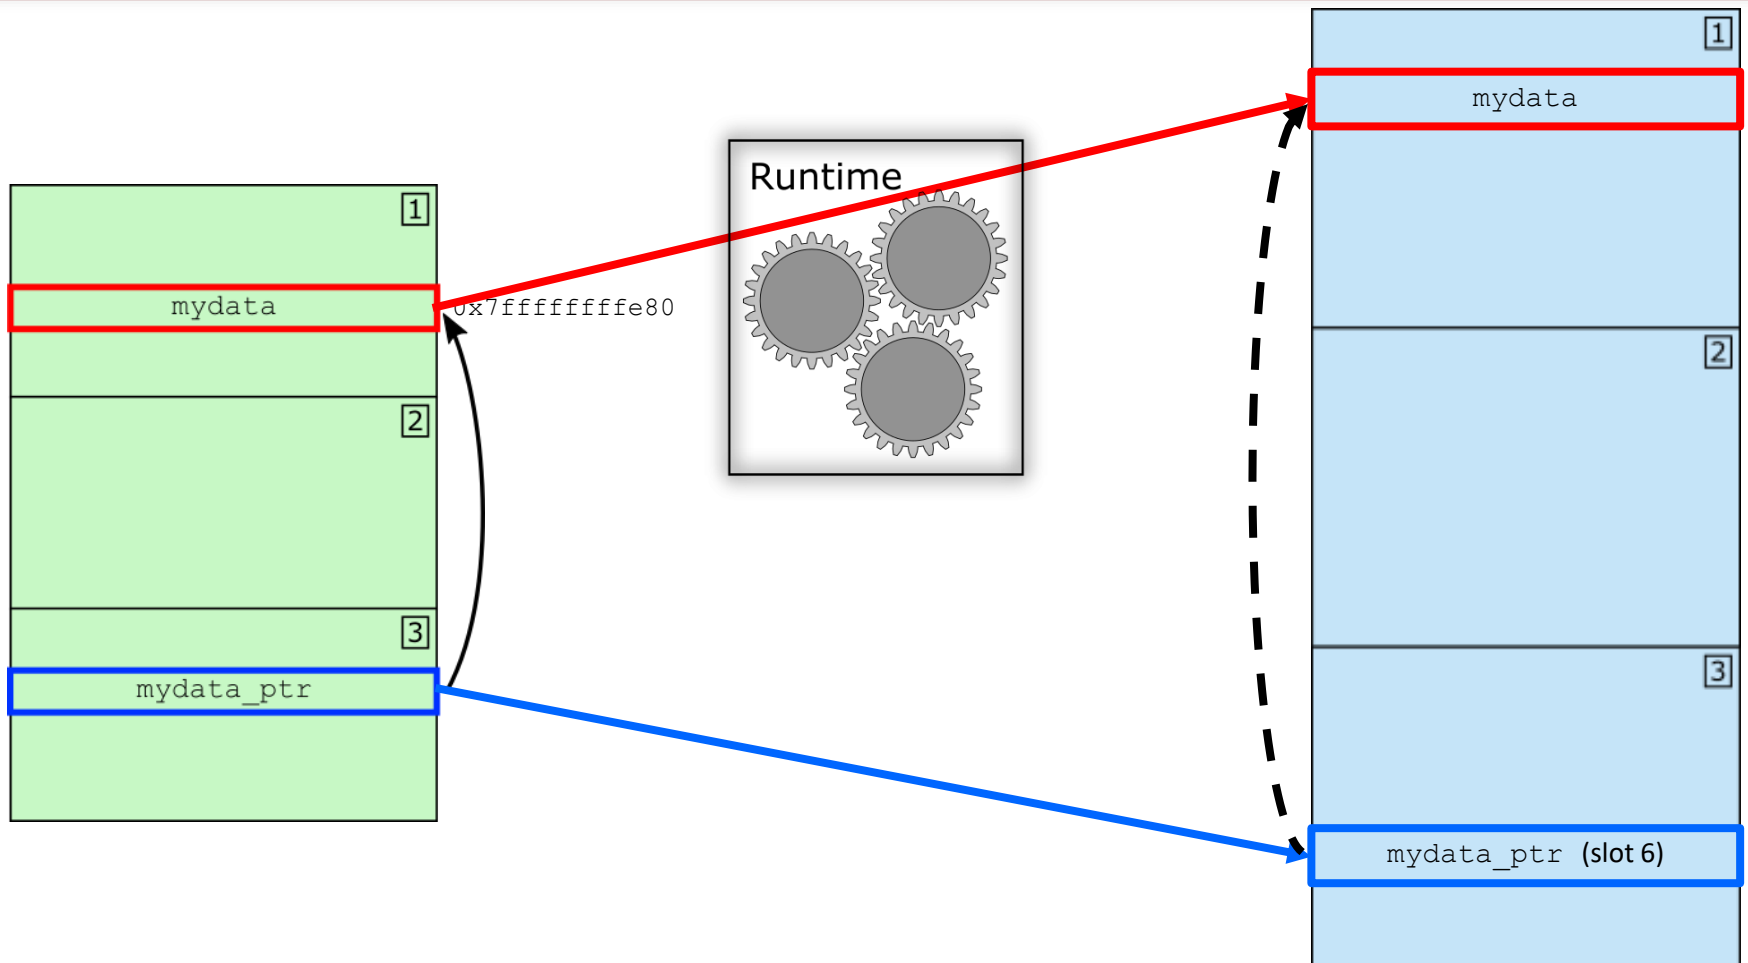




# State Transformation



# State Transformation



# Runtime Summary

---

- Migrations are ***proposed*** by threads/applications
- Threads check to see if migration is proposed at migration points
- If so, invoke migration procedure
  - Transform thread's register set, stack between target-specific formats
  - Invoke kernel's thread migration service
  - Bootstrap on destination & return to normal execution